Microsoft

T-SQL Fundamentals

Third Edition

Professional

Itzik Ben-Gan
SolidQ

FREE SAMPLE CHAPTER

SHARE WITH OTHERS

Microsoft

T-SQL Fundamentals, Third Edition

Itzik Ben-Gan

PUBLISHED BY
Microsoft Press
A division of Microsoft Corporation
One Microsoft Way
Redmond, Washington 98052-6399

Copyright © 2016 by Itzik Ben-Gan

All rights reserved. No part of the contents of this book may be reproduced or transmitted in any form or by any means without the written permission of the publisher.

Library of Congress Control Number: 2015955815
ISBN: 978-1-5093-0200-0

Printed and bound in the United States of America.

First Printing

Microsoft Press books are available through booksellers and distributors worldwide. If you need support related to this book, email Microsoft Press Support at mspinput@microsoft.com. Please tell us what you think of this book at http://aka.ms/tellpress.

This book is provided "as-is" and expresses the author's views and opinions. The views, opinions and information expressed in this book, including URL and other Internet website references, may change without notice.

Some examples depicted herein are provided for illustration only and are fictitious. No real association or connection is intended or should be inferred.

Microsoft and the trademarks listed at http://www.microsoft.com on the "Trademarks" webpage are trademarks of the Microsoft group of companies. All other marks are property of their respective owners.

Acquisitions Editor: Devon Musgrave
Developmental Editor: Carol Dillingham
Project Editor: Carol Dillingham
Editorial Production: Christian Holdener; S4Carlisle Publishing Services
Technical Reviewer: Bob Beauchemin; Technical Review services provided by Content Master, a member of CM Group, Ltd.
Copyeditor: Roger Leblanc
Indexer: Maureen Johnson, MoJo's Indexing Services
Cover: Twist Creative • Seattle and Joel Panchot

To Dato,
To live in hearts we leave behind,
Is not to die.

—Thomas Campbell

This page intentionally left blank

Contents at a glance

	Introduction	*xix*
CHAPTER 1	**Background to T-SQL querying and programming**	**1**
CHAPTER 2	**Single-table queries**	**27**
CHAPTER 3	**Joins**	**103**
CHAPTER 4	**Subqueries**	**133**
CHAPTER 5	**Table expressions**	**161**
CHAPTER 6	**Set operators**	**193**
CHAPTER 7	**Beyond the fundamentals of querying**	**213**
CHAPTER 8	**Data modification**	**249**
CHAPTER 9	**Temporal tables**	**297**
CHAPTER 10	**Transactions and concurrency**	**319**
CHAPTER 11	**Programmable objects**	**361**
	Appendix: Getting started	*395*
	Index	*415*

This page intentionally left blank

Contents

Introduction ..... *xix*

Chapter 1 Background to T-SQL querying and programming 1

Theoretical background 1
- SQL 2
- Set theory 3
- Predicate logic 4
- The relational model 4
- Types of database systems 10

SQL Server architecture 12
- The ABCs of Microsoft RDBMS flavors 12
- SQL Server instances 14
- Databases 15
- Schemas and objects 18

Creating tables and defining data integrity 19
- Creating tables 20
- Defining data integrity 22

Conclusion 25

Chapter 2 Single-table queries 27

Elements of the *SELECT* statement 27
- The *FROM* clause 29
- The *WHERE* clause 31
- The *GROUP BY* clause 32
- The *HAVING* clause 35
- The *SELECT* clause 36

What do you think of this book? We want to hear from you!

Microsoft is interested in hearing your feedback so we can improve our books and learning resources for you. To participate in a brief survey, please visit:

http://aka.ms/tellpress

The *ORDER BY* clause 41
The *TOP* and *OFFSET-FETCH* filters 43
A quick look at window functions 47
Predicates and operators 49
CASE expressions 52
NULLs 54
All-at-once operations 58
Working with character data 60
Data types 60
Collation 61
Operators and functions 63
The *LIKE* predicate 71
Working with date and time data 73
Date and time data types 74
Literals 74
Working with date and time separately 78
Filtering date ranges 79
Date and time functions 80
Querying metadata 89
Catalog views 90
Information schema views 91
System stored procedures and functions 91
Conclusion 93
Exercises 93
Exercise 1 93
Exercise 2 93
Exercise 3 94
Exercise 4 94
Exercise 5 95
Exercise 6 95
Exercise 7 95
Exercise 8 96
Exercise 9 96
Exercise 10 97

Solutions 97
Exercise 1 98
Exercise 2 98
Exercise 3 99
Exercise 4 99
Exercise 5 99
Exercise 6 100
Exercise 7 101
Exercise 8 102
Exercise 9 102
Exercise 10 102

Chapter 3 Joins 103

Cross joins 103
ISO/ANSI SQL-92 syntax 104
ISO/ANSI SQL-89 syntax 105
Self cross joins 105
Producing tables of numbers 106
Inner joins 107
ISO/ANSI SQL-92 syntax 107
ISO/ANSI SQL-89 syntax 108
Inner join safety 109
More join examples 109
Composite joins 110
Non-equi joins 111
Multi-join queries 113
Outer joins 113
Fundamentals of outer joins 114
Beyond the fundamentals of outer joins 116
Conclusion 123
Exercises 123
Exercise 1-1 123
Exercise 1-2 (optional, advanced) 124
Exercise 2 125

Exercise 3 . . . 125
Exercise 4 . . . 126
Exercise 5 . . . 126
Exercise 6 . . . 127
Exercise 7 (optional, advanced) . . . 127
Exercise 8 (optional, advanced) . . . 128
Exercise 9 (optional, advanced) . . . 128

Solutions . . . 129
Exercise 1-1 . . . 129
Exercise 1-2 . . . 129
Exercise 2 . . . 129
Exercise 3 . . . 130
Exercise 4 . . . 130
Exercise 5 . . . 131
Exercise 6 . . . 131
Exercise 7 . . . 131
Exercise 8 . . . 132
Exercise 9 . . . 132

Chapter 4 Subqueries 133

Self-contained subqueries . . . 133
Self-contained scalar subquery examples . . . 134
Self-contained multivalued subquery examples . . . 136

Correlated subqueries . . . 139
The *EXISTS* predicate . . . 141

Beyond the fundamentals of subqueries . . . 143
Returning previous or next values . . . 143
Using running aggregates . . . 144
Dealing with misbehaving subqueries . . . 145

Conclusion . . . 149

Exercises . . . 150
Exercise 1 . . . 150
Exercise 2 (optional, advanced) . . . 150
Exercise 3 . . . 151

Exercise 4 151
Exercise 5 152
Exercise 6 152
Exercise 7 (optional, advanced) 153
Exercise 8 (optional, advanced) 153
Exercise 9 154
Exercise 10 (optional, advanced) 154
Solutions 154
Exercise 1 154
Exercise 2 155
Exercise 3 155
Exercise 4 155
Exercise 5 156
Exercise 6 156
Exercise 7 156
Exercise 8 157
Exercise 9 157
Exercise 10 158

Chapter 5 Table expressions 161

Derived tables 161
Assigning column aliases 163
Using arguments 164
Nesting 165
Multiple references 166
Common table expressions 167
Assigning column aliases in CTEs 167
Using arguments in CTEs 168
Defining multiple CTEs 168
Multiple references in CTEs 169
Recursive CTEs 169
Views 171
Views and the *ORDER BY* clause 172
View options 175

Inline table-valued functions . . . 178
The *APPLY* operator . . . 180
Conclusion . . . 183
Exercises . . . 183
Exercise 1 . . . 183
Exercise 2-1 . . . 184
Exercise 2-2 . . . 184
Exercise 3-1 . . . 184
Exercise 3-2 . . . 185
Exercise 4 (optional, advanced) . . . 185
Exercise 5-1 . . . 186
Exercise 5-2 (optional, advanced) . . . 187
Exercise 6-1 . . . 187
Exercise 6-2 . . . 188
Solutions . . . 188
Exercise 1 . . . 188
Exercise 2-1 . . . 189
Exercise 2-2 . . . 189
Exercise 3-1 . . . 189
Exercise 3-2 . . . 189
Exercise 4 . . . 190
Exercise 5-1 . . . 190
Exercise 5-2 . . . 191
Exercise 6-1 . . . 191
Exercise 6-2 . . . 192

Chapter 6 Set operators 193

The *UNION* operator . . . 194
The *UNION ALL* operator . . . 194
The *UNION* (*DISTINCT*) operator . . . 195
The *INTERSECT* Operator . . . 196
The *INTERSECT* (*DISTINCT*) operator . . . 196
The *INTERSECT ALL* operator . . . 197
The *EXCEPT* operator . . . 199

The *EXCEPT* (*DISTINCT*) operator 199
The *EXCEPT ALL* operator 200
Precedence 201
Circumventing unsupported logical phases 203
Conclusion 204
Exercises 204
Exercise 1 204
Exercise 2 205
Exercise 3 205
Exercise 4 206
Exercise 5 207
Exercise 6 (optional, advanced) 207
Solutions 208
Exercise 1 208
Exercise 2 208
Exercise 3 209
Exercise 4 210
Exercise 5 210
Exercise 6 211

Chapter 7 Beyond the fundamentals of querying 213

Window functions 213
Ranking window functions 216
Offset window functions 219
Aggregate window functions 221
Pivoting data 224
Pivoting with a grouped query 226
Pivoting with the *PIVOT* operator 227
Unpivoting data 230
Unpivoting with the *APPLY* operator 231
Unpivoting with the *UNPIVOT* operator 233
Grouping sets 234
The *GROUPING SETS* subclause 236
The *CUBE* subclause 236

The *ROLLUP* subclause 236
The *GROUPING* and *GROUPING_ID* functions 238
Conclusion 240
Exercises 241
Exercise 1 241
Exercise 2 241
Exercise 3 242
Exercise 4 242
Exercise 5 243
Exercise 6 244
Solutions 245
Exercise 1 245
Exercise 2 245
Exercise 3 245
Exercise 4 246
Exercise 5 247
Exercise 6 247

Chapter 8 Data modification 249

Inserting data 249
The *INSERT VALUES* statement 249
The *INSERT SELECT* statement 251
The *INSERT EXEC* statement 251
The *SELECT INTO* statement 252
The *BULK INSERT* statement 253
The identity property and the sequence object 254
Deleting data 262
The *DELETE* statement 263
The *TRUNCATE* statement 264
DELETE based on a join 265
Updating data 266
The *UPDATE* statement 267
UPDATE based on a join 268
Assignment *UPDATE* 270

Merging data 271
Modifying data through table expressions 276
Modifications with *TOP* and *OFFSET-FETCH* 278
The *OUTPUT* clause 280
INSERT with *OUTPUT* 280
DELETE with *OUTPUT* 282
UPDATE with *OUTPUT* 283
MERGE with *OUTPUT* 284
Nested DML 285
Conclusion 287
Exercises 287
Exercise 1 287
Exercise 2 288
Exercise 3 288
Exercise 4 288
Exercise 5 290
Exercise 6 290
Solutions 291
Exercise 1 291
Exercise 2 293
Exercise 3 293
Exercise 4 294
Exercise 5 294
Exercise 6 295

Chapter 9 Temporal tables 297

Creating tables 298
Modifying data 301
Querying data 304
Conclusion 311
Exercises 311
Exercise 1 311
Exercise 2 311

Exercise 3 . . . 312
Exercise 4 . . . 313
Solutions . . . 313
Exercise 1 . . . 313
Exercise 2 . . . 314
Exercise 3 . . . 316
Exercise 4 . . . 317

Chapter 10 Transactions and concurrency 319

Transactions . . . 319
Locks and blocking . . . 322
Locks . . . 322
Troubleshooting blocking . . . 325
Isolation levels . . . 332
The *READ UNCOMMITTED* isolation level . . . 333
The *READ COMMITTED* isolation level . . . 334
The *REPEATABLE READ* isolation level . . . 335
The *SERIALIZABLE* isolation level . . . 337
Isolation levels based on row versioning . . . 339
Summary of isolation levels . . . 345
Deadlocks . . . 345
Conclusion . . . 348
Exercises . . . 348
Exercise 1-1 . . . 349
Exercise 1-2 . . . 349
Exercise 1-3 . . . 349
Exercise 1-4 . . . 349
Exercise 1-5 . . . 350
Exercise 1-6 . . . 350
Exercise 2-1 . . . 350
Exercise 2-2 . . . 351
Exercise 2-3 . . . 352
Exercise 2-4 . . . 353
Exercise 2-5 . . . 354

Exercise 2-6 356
Exercise 3-1 357
Exercise 3-2 358
Exercise 3-3 358
Exercise 3-4 358
Exercise 3-5 358
Exercise 3-6 358
Exercise 3-7 359

Chapter 11 Programmable objects 361

Variables 361
Batches 363
A batch as a unit of parsing 364
Batches and variables 365
Statements that cannot be combined in the same batch 365
A batch as a unit of resolution 366
The *GO n* option 366
Flow elements 367
The *IF . . . ELSE* flow element 367
The *WHILE* flow element 368
Cursors 370
Temporary tables 374
Local temporary tables 374
Global temporary tables 376
Table variables 377
Table types 378
Dynamic SQL 379
The *EXEC* command 380
The *sp_executesql* stored procedure 380
Using *PIVOT* with Dynamic SQL 381
Routines 383
User-defined functions 383
Stored procedures 385
Triggers 387

Error handling 390
Conclusion 394

Appendix: Getting started *395*

Index *415*

What do you think of this book? We want to hear from you!

Microsoft is interested in hearing your feedback so we can improve our books and learning resources for you. To participate in a brief survey, please visit:

http://aka.ms/tellpress

Introduction

This book walks you through your first steps in T-SQL (also known as *Transact-SQL*), which is the Microsoft SQL Server dialect of the ISO and ANSI standards for SQL. You'll learn the theory behind T-SQL querying and programming and how to develop T-SQL code to query and modify data, and you'll get an overview of programmable objects.

Although this book is intended for beginners, it's not merely a set of procedures for readers to follow. It goes beyond the syntactical elements of T-SQL and explains the logic behind the language and its elements.

Occasionally, the book covers subjects that might be considered advanced for readers who are new to T-SQL; therefore, you should consider those sections to be optional reading. The text will indicate when a section is considered more advanced and is provided as optional reading. If you feel comfortable with the material discussed in the book up to that point, you might want to tackle these more advanced subjects; otherwise, feel free to skip those sections and return to them after you gain more experience.

Many aspects of SQL are unique to the language and very different from other programming languages. This book helps you adopt the right state of mind and gain a true understanding of the language elements. You learn how to think in relational terms and follow good SQL programming practices.

The book is not version specific; it does, however, cover language elements that were introduced in recent versions of SQL Server, including SQL Server 2016. When I discuss language elements that were introduced recently, I specify the version in which they were added.

Besides being available as a box product, SQL Server is also available as a cloud-based service called Microsoft Azure SQL Database, or in short, just SQL Database. The code samples in this book were tested against both a box SQL Server product and Azure SQL Database. The book's companion content (available at *http://aka.ms/T-SQLFund3e/downloads*) provides information about compatibility issues between the flavors.

To complement the learning experience, the book provides exercises you can use to practice what you learn. The book occasionally provides optional exercises that are more advanced. Those exercises are intended for readers who feel comfortable with the material and want to challenge themselves with more difficult problems. The optional exercises for advanced readers are labeled as such.

Who should read this book

This book is intended for T-SQL developers, database administrators (DBAs), business intelligence (BI) practitioners, data scientists, report writers, analysts, architects, and SQL Server power users who just started working with SQL Server and who need to write queries and develop code using Transact-SQL.

Assumptions

To get the most out of this book, you should have working experience with Microsoft Windows and with applications based on Windows. You should also be familiar with basic concepts of relational database management systems.

This book might not be for you if...

Not every book is aimed at every possible audience. This book covers fundamentals. It's mainly aimed at T-SQL practitioners with little or no experience. This book might not be for you if you're an advanced T-SQL practitioner with many years of experience. With that said, several readers of the previous editions of this book have mentioned that—even though they already had years of experience—they still found the book useful for filling gaps in their knowledge.

Organization of this book

This book starts with a theoretical background to T-SQL querying and programming in Chapter 1, laying the foundations for the rest of the book, and basic coverage of creating tables and defining data integrity. The book moves on to various aspects of querying and modifying data in Chapters 2 through 9, and then moves to a discussion of concurrency and transactions in Chapter 10. Finally, it provides an overview of programmable objects in Chapter 11.

Here's a list of the chapters along with a short description of the content in each chapter:

- Chapter 1, "Background to T-SQL querying and programming," provides the theoretical background for SQL, set theory, and predicate logic. It examines the relational model, describes SQL Server's architecture, and explains how to create tables and define data integrity.

- Chapter 2, "Single-table queries," covers various aspects of querying a single table by using the *SELECT* statement.
- Chapter 3, "Joins," covers querying multiple tables by using joins, including cross joins, inner joins, and outer joins.
- Chapter 4, "Subqueries," covers queries within queries, otherwise known as *subqueries*.
- Chapter 5, "Table expressions," covers derived tables, Common Table Expressions (CTEs), views, inline table-valued functions, and the *APPLY* operator.
- Chapter 6, "Set operators," covers the set operators *UNION*, *INTERSECT*, and *EXCEPT*.
- Chapter 7, "Beyond the fundamentals of querying," covers window functions, pivoting, unpivoting, and working with grouping sets.
- Chapter 8, "Data modification," covers inserting, updating, deleting, and merging data.
- Chapter 9, "Temporal tables," covers system-versioned temporal tables.
- Chapter 10, "Transactions and concurrency," covers concurrency of user connections that work with the same data simultaneously; it covers transactions, locks, blocking, isolation levels, and deadlocks.
- Chapter 11, "Programmable objects," provides an overview of the T-SQL programming capabilities in SQL Server.
- The book also provides an appendix, "Getting started," to help you set up your environment, download the book's source code, install the TSQLV4 sample database, start writing code against SQL Server, and learn how to get help by working with SQL Server Books Online.

System requirements

The Appendix, "Getting started," explains which editions of SQL Server 2016 you can use to work with the code samples included with this book. Each edition of SQL Server might have different hardware and software requirements, and those requirements are well documented in SQL Server Books Online under "Hardware and Software Requirements for Installing SQL Server 2016" at the following URL: *https://msdn.microsoft.com/en-us/library/ms143506.aspx*. The Appendix also explains how to work with SQL Server Books Online.

If you're connecting to Azure SQL Database, hardware and server software are handled by Microsoft, so those requirements are irrelevant in this case.

You will need SQL Server Management Studio (SSMS) to run the code samples against both SQL Server 2016 and Azure SQL Database. You can download SSMS and find information about the supported operating systems at the following URL: *https://msdn.microsoft.com/en-us/library/mt238290.aspx*.

Installing and using the source code

Most of the chapters in this book include exercises that let you interactively try out new material learned in the main text. All source code, including exercises and solutions, can be downloaded from the following webpage:

http://aka.ms/T-SQLFund3e/downloads

Follow the instructions to download the TSQLFundamentalsYYYYMMDD.zip file, where YYYYMMDD reflects the last update date of the source code.

Refer to the Appendix, "Getting started," for details about the source code.

Acknowledgments

A number of people contributed to making this book a reality, either directly or indirectly, and deserve thanks and recognition. It's certainly possible I omitted some names unintentionally, and I apologize for this ahead of time.

To Lilach: You're the one who makes me want to be good at what I do. Besides being my inspiration in life, you always take an active role in my books, helping to review the text for the first time.

To my mom, Mila, and to my siblings, Mickey and Ina: Thank you for the constant support and for accepting the fact that I'm away. To my dad, Gabi, who loved puzzles, logic, and numbers: I attribute my affinity to SQL to you; we all miss you a lot.

To the technical reviewer of the book, Bob Beauchemin: You've been around in the SQL Server community for many years. I was always impressed by your extensive knowledge and was happy that you agreed to work with me on this book.

To Steve Kass, Dejan Sarka, Gianluca Hotz, and Herbert Albert: Thanks for your valuable advice during the planning phases of the book. I had to make some hard decisions in terms of what to include and what not to include in the book, and your advice was very helpful.

Many thanks to the book's editors. To Devon Musgrave, who played the acquisitions editor role: You are the one who made this book a reality and handled all the initial

stages. I realize that this book is likely one of many you were responsible for, and I'd like to thank you for dedicating the time and effort that you did. To Carol Dillingham, the book's developmental editor and project editor, many thanks for your excellent handling; I always enjoy working with you on my books. Also thanks to Roger LeBlanc for his fine copy edits, and to Christian Holdener for his project management.

To SolidQ, my company for over a decade: It's gratifying to be part of such a great company that evolved into what it is today. The members of this company are much more than colleagues to me; they are partners, friends, and family. To Fernando G. Guerrero, Antonio Soto, and Douglas McDowell: thanks for leading the company. To my many colleagues: It's an honor to be part of this amazingly talented group.

To members of the Microsoft SQL Server development team, past and present: Tobias Ternstrom, Lubor Kollar, Umachandar Jayachandran (UC), Boris Baryshnikov, Conor Cunningham, Kevin Farlee, Josde Bruijn, Marc Friedman, Drazen Sumic, Borko Novakovic, Milan Stojic, Milan Ruzic, Jovan Popovic, Lindsey Allen, Craig Freedman, Campbell Fraser, Eric Hanson, Mark Souza, Dave Campbell, César Galindo-Legaria, Pedro Lopes, and I'm sure many others. Thanks for creating such a great product, and thanks for all the time you spent meeting with me and responding to my emails, addressing my questions, and answering my requests for clarification.

To members of the SQL Server Pro editorial team, Tim Ford and Debra Donston-Miller: I've been writing for the magazine for almost two decades, and I'm grateful for the opportunity to share my knowledge with the magazine's readers.

To Data Platform MVPs, past and present: Paul White, Alejandro Mesa, Erland Sommarskog, Aaron Bertrand, Tibor Karaszi, Benjamin Nevarez, Simon Sabin, Darren Green, Allan Mitchell, Tony Rogerson, and many others—and to the Data Platform MVP lead, Jennifer Moser. This is a great program that I'm grateful for and proud to be part of. The level of expertise of this group is amazing, and I'm always excited when we all get to meet, both to share ideas and just to catch up at a personal level over beers.

Finally, to my students: Teaching about T-SQL is what drives me. It's my passion. Thanks for allowing me to fulfill my calling and for all the great questions that make me seek more knowledge.

Errata, updates, & book support

We've made every effort to ensure the accuracy of this book and its companion content. You can access updates to this book—in the form of a list of submitted errata and their related corrections—at:

http://aka.ms/T-SQLFund3e/errata

If you discover an error that is not already listed, please submit it to us at the same page.

If you need additional support, email Microsoft Press Book Support at *mspinput@microsoft.com*.

Please note that product support for Microsoft software and hardware is not offered through the previous addresses. For help with Microsoft software or hardware, go to *http://support.microsoft.com*.

Free ebooks from Microsoft Press

From technical overviews to in-depth information on special topics, the free ebooks from Microsoft Press cover a wide range of topics. These ebooks are available in PDF, EPUB, and Mobi for Kindle formats, ready for you to download at:

http://aka.ms/mspressfree

Check back often to see what is new!

We want to hear from you

At Microsoft Press, your satisfaction is our top priority, and your feedback our most valuable asset. Please tell us what you think of this book at:

http://aka.ms/tellpress

We know you're busy, so we've kept it short with just a few questions. Your answers go directly to the editors at Microsoft Press. (No personal information will be requested.) Thanks in advance for your input!

Stay in touch

Let's keep the conversation going! We're on Twitter: *http://twitter.com/MicrosoftPress*.

CHAPTER 1

Background to T-SQL querying and programming

You're about to embark on a journey to a land that is like no other—a land that has its own set of laws. If reading this book is your first step in learning Transact-SQL (T-SQL), you should feel like Alice—just before she started her adventures in Wonderland. For me, the journey has not ended; instead, it's an ongoing path filled with new discoveries. I envy you; some of the most exciting discoveries are still ahead of you!

I've been involved with T-SQL for many years: teaching, speaking, writing, and consulting about it. For me, T-SQL is more than just a language—it's a way of thinking. In my first few books about T-SQL, I've written extensively on advanced topics, and for years, I have postponed writing about fundamentals. This is not because T-SQL fundamentals are simple or easy—in fact, it's just the opposite. The apparent simplicity of the language is misleading. I could explain the language syntax elements in a superficial manner and have you writing queries within minutes. But that approach would only hold you back in the long run and make it harder for you to understand the essence of the language.

Acting as your guide while you take your first steps in this realm is a big responsibility. I wanted to make sure that I spent enough time and effort exploring and understanding the language before writing about fundamentals. T-SQL is deep; learning the fundamentals the right way involves much more than just understanding the syntax elements and coding a query that returns the right output. You pretty much need to forget what you know about other programming languages and start thinking in terms of T-SQL.

Theoretical background

SQL stands for *Structured Query Language*. SQL is a standard language that was designed to query and manage data in relational database management systems (RDBMSs). An RDBMS is a database management system based on the relational model (a semantic model for representing data), which in turn is based on two mathematical branches: set theory and predicate logic. Many other programming languages and various aspects of computing evolved pretty much as a result of intuition. In contrast, to the degree that SQL is based on the relational model, it is based on a firm foundation—applied mathematics. T-SQL thus sits on wide and solid shoulders. Microsoft provides T-SQL as a dialect of, or extension to, SQL in Microsoft SQL Server data-management software, its RDBMS.

This section provides a brief theoretical background about SQL, set theory and predicate logic, the relational model, and types of database systems. Because this book is neither a mathematics book nor a design/data-modeling book, the theoretical information provided here is informal and by no means complete. The goals are to give you a context for the T-SQL language and to deliver the key points that are integral to correctly understanding T-SQL later in the book.

Language independence

The relational model is language independent. That is, you can apply data management and manipulation following the relational model's principles with languages other than SQL—for example, with C# in an object model. Today it is common to see RDBMSs that support languages other than a dialect of SQL, such as the CLR integration in SQL Server, with which you can handle tasks that historically you handled mainly with SQL, such as data manipulation.

Also, you should realize from the start that SQL deviates from the relational model in several ways. Some even say that a new language—one that more closely follows the relational model—should replace SQL. But to date, SQL is the industrial language used by all leading RDBMSs in practice.

See Also *For details about the deviations of SQL from the relational model, as well as how to use SQL in a relational way, see this book on the topic:* SQL and Relational Theory: How to Write Accurate SQL Code, Third Edition *by C. J. Date (O'Reilly Media, 2015).*

SQL

SQL is both an ANSI and ISO standard language based on the relational model, designed for querying and managing data in an RDBMS.

In the early 1970s, IBM developed a language called SEQUEL (short for Structured English QUEry Language) for its RDBMS product called System R. The name of the language was later changed from SEQUEL to SQL because of a trademark dispute. SQL first became an ANSI standard in 1986, and then an ISO standard in 1987. Since 1986, the American National Standards Institute (ANSI) and the International Organization for Standardization (ISO) have been releasing revisions for the SQL standard every few years. So far, the following standards have been released: SQL-86 (1986), SQL-89 (1989), SQL-92 (1992), SQL:1999 (1999), SQL:2003 (2003), SQL:2006 (2006), SQL:2008 (2008), and SQL:2011 (2011). The SQL standard is made of multiple parts. Part 1 (Framework) and Part 2 (Foundation) pertain to the SQL language, whereas the other parts define standard extensions, such as SQL for XML and SQL-Java integration.

Interestingly, SQL resembles English and is also very logical. Unlike many programming languages, which use an imperative programming paradigm, SQL uses a declarative one. That is, SQL requires you to specify *what* you want to get and not *how* to get it, letting the RDBMS figure out the physical mechanics required to process your request.

SQL has several categories of statements, including Data Definition Language (DDL), Data Manipulation Language (DML), and Data Control Language (DCL). DDL deals with object definitions

and includes statements such as *CREATE*, *ALTER*, and *DROP*. DML allows you to query and modify data and includes statements such as *SELECT*, *INSERT*, *UPDATE*, *DELETE*, *TRUNCATE*, and *MERGE*. It's a common misunderstanding that DML includes only data-modification statements, but as I mentioned, it also includes *SELECT*. Another common misunderstanding is that *TRUNCATE* is a DDL statement, but in fact it is a DML statement. DCL deals with permissions and includes statements such as *GRANT* and *REVOKE*. This book focuses on DML.

T-SQL is based on standard SQL, but it also provides some nonstandard/proprietary extensions. Moreover, T-SQL does not implement all of standard SQL. In other words, T-SQL is both a subset and a superset of SQL. When describing a language element for the first time, I'll typically mention whether it is standard.

Set theory

Set theory, which originated with the mathematician Georg Cantor, is one of the mathematical branches on which the relational model is based. Cantor's definition of a set follows:

> *By a "set" we mean any collection M into a whole of definite, distinct objects m (which are called the "elements" of M) of our perception or of our thought.*
>
> *—Joseph W. Dauben and Georg Cantor (Princeton University Press, 1990)*

Every word in the definition has a deep and crucial meaning. The definitions of a set and set membership are axioms that are not supported by proofs. Each element belongs to a universe, and either is or is not a member of the set.

Let's start with the word *whole* in Cantor's definition. A set should be considered a single entity. Your focus should be on the collection of objects as opposed to the individual objects that make up the collection. Later on, when you write T-SQL queries against tables in a database (such as a table of employees), you should think of the set of employees as a whole rather than the individual employees. This might sound trivial and simple enough, but apparently many programmers have difficulty adopting this way of thinking.

The word *distinct* means that every element of a set must be unique. Jumping ahead to tables in a database, you can enforce the uniqueness of rows in a table by defining key constraints. Without a key, you won't be able to uniquely identify rows, and therefore the table won't qualify as a set. Rather, the table would be a *multiset* or a *bag*.

The phrase *of our perception or of our thought* implies that the definition of a set is subjective. Consider a classroom: one person might perceive a set of people, whereas another might perceive a set of students and a set of teachers. Therefore, you have a substantial amount of freedom in defining sets. When you design a data model for your database, the design process should carefully consider the subjective needs of the application to determine adequate definitions for the entities involved.

As for the word *object*, the definition of a set is not restricted to physical objects, such as cars or employees, but rather is relevant to abstract objects as well, such as prime numbers or lines.

What Cantor's definition of a set leaves out is probably as important as what it includes. Notice that the definition doesn't mention any order among the set elements. The order in which set elements are listed is not important. The formal notation for listing set elements uses curly brackets: {a, b, c}. Because order has no relevance, you can express the same set as *{b, a, c}* or *{b, c, a}*. Jumping ahead to the set of attributes (called *columns* in SQL) that make up the heading of a relation (called a *table* in SQL), an element is supposed to be identified by name—not by ordinal position.

Similarly, consider the set of tuples (called *rows* by SQL) that make up the body of the relation; an element is identified by its key values—not by position. Many programmers have a hard time adapting to the idea that, with respect to querying tables, there is no order among the rows. In other words, a query against a table can return table rows in *any order* unless you explicitly request that the data be sorted in a specific way, perhaps for presentation purposes.

Predicate logic

Predicate logic, whose roots reach back to ancient Greece, is another branch of mathematics on which the relational model is based. Dr. Edgar F. Codd, in creating the relational model, had the insight to connect predicate logic to both the management and querying of data. Loosely speaking, a *predicate* is a property or an expression that either holds or doesn't hold—in other words, is either true or false. The relational model relies on predicates to maintain the logical integrity of the data and define its structure. One example of a predicate used to enforce integrity is a constraint defined in a table called *Employees* that allows only employees with a salary greater than zero to be stored in the table. The predicate is "salary greater than zero" (T-SQL expression: *salary > 0*).

You can also use predicates when filtering data to define subsets, and more. For example, if you need to query the *Employees* table and return only rows for employees from the sales department, you use the predicate "department equals sales" in your query filter (T-SQL expression: *department = 'sales'*).

In set theory, you can use predicates to define sets. This is helpful because you can't always define a set by listing all its elements (for example, infinite sets), and sometimes for brevity it's more convenient to define a set based on a property. As an example of an infinite set defined with a predicate, the set of all prime numbers can be defined with the following predicate: "*x* is a positive integer greater than 1 that is divisible only by 1 and itself." For any specified value, the predicate is either true or not true. The set of all prime numbers is the set of all elements for which the predicate is true. As an example of a finite set defined with a predicate, the set *{0, 1, 2, 3, 4, 5, 6, 7, 8, 9}* can be defined as the set of all elements for which the following predicate holds true: "*x* is an integer greater than or equal to 0 and smaller than or equal to 9."

The relational model

The relational model is a semantic model for data management and manipulation and is based on set theory and predicate logic. As mentioned earlier, it was created by Dr. Edgar F. Codd, and later explained and developed by Chris Date, Hugh Darwen, and others. The first version of the relational model was proposed by Codd in 1969 in an IBM research report called "Derivability, Redundancy, and

Consistency of Relations Stored in Large Data Banks." A revised version was proposed by Codd in 1970 in a paper called "A Relational Model of Data for Large Shared Data Banks," published in the journal *Communications of the ACM*.

The goal of the relational model is to enable consistent representation of data with minimal or no redundancy and without sacrificing completeness, and to define data integrity (enforcement of data consistency) as part of the model. An RDBMS is supposed to implement the relational model and provide the means to store, manage, enforce the integrity of, and query data. The fact that the relational model is based on a strong mathematical foundation means that given a certain data-model instance (from which a physical database will later be generated), you can tell with certainty when a design is flawed, rather than relying solely on intuition.

The relational model involves concepts such as propositions, predicates, relations, tuples, attributes, and more. For nonmathematicians, these concepts can be quite intimidating. The sections that follow cover some key aspects of the model in an informal, nonmathematical manner and explain how they relate to databases.

Propositions, predicates, and relations

The common belief that the term *relational* stems from relationships between tables is incorrect. "Relational" actually pertains to the mathematical term *relation*. In set theory, a relation is a representation of a set. In the relational model, a relation is a set of related information, with the counterpart in SQL being a table—albeit not an exact counterpart. A key point in the relational model is that a single relation should represent a single set (for example, *Customers*). Note that operations on relations (based on relational algebra) result in a relation (for example, a join between two relations).

Note The relational model distinguishes between a *relation* and a *relation variable*, but to keep things simple, I won't get into this distinction. Instead, I'll use the term *relation* for both cases. Also, a relation is made of a heading and a body. The heading consists of a set of attributes (called *columns* in SQL), where each element is identified by an attribute name and a type name. The body consists of a set of tuples (called *rows* in SQL), where each element is identified by a key. To keep things simple, I'll refer to a table as a set of rows.

When you design a data model for a database, you represent all data with relations (tables). You start by identifying propositions that you will need to represent in your database. A proposition is an assertion or a statement that must be true or false. For example, the statement, "Employee Itzik Ben-Gan was born on February 12, 1971, and works in the IT department" is a proposition. If this proposition is true, it will manifest itself as a row in a table of *Employees*. A false proposition simply won't manifest itself. This presumption is known as the *closed-world assumption (CWA)*.

The next step is to formalize the propositions. You do this by taking out the actual data (the body of the relation) and defining the structure (the heading of the relation)—for example, by creating predicates out of propositions. You can think of predicates as parameterized propositions. The heading of a relation comprises a set of attributes. Note the use of the term "set"; in the relational model,

attributes are unordered and distinct. An attribute is identified by an attribute name and a type name. For example, the heading of an *Employees* relation might consist of the following attributes (expressed as pairs of attribute names and type names): *employeeid* integer, *firstname* character string, *lastname* character string, *birthdate* date, and *departmentid* integer.

A *type* is one of the most fundamental building blocks for relations. A type constrains an attribute to a certain set of possible or valid values. For example, the type *INT* is the set of all integers in the range –2,147,483,648 to 2,147,483,647. A type is one of the simplest forms of a predicate in a database because it restricts the attribute values that are allowed. For example, the database would not accept a proposition where an employee birth date is February 31, 1971 (not to mention a birth date stated as something like "abc!"). Note that types are not restricted to base types such as integers or character strings; a type also can be an enumeration of possible values, such as an enumeration of possible job positions. A type can be simple or complex. Probably the best way to think of a type is as a class—encapsulated data and the behavior supporting it. An example of a complex type is a geometry type that supports polygons.

Missing values

One aspect of the relational model is the source of many passionate debates—whether predicates should be restricted to two-valued logic. That is, in two-valued predicate logic, a predicate is either true or false. If a predicate is not true, it must be false. Use of two-valued predicate logic follows a mathematical law called "the law of excluded middle." However, some say that there's room for three-valued (or even four-valued) predicate logic, taking into account cases where values are missing. A predicate involving a missing value yields neither *true* nor *false*—it yields *unknown*.

Take, for example, a mobile phone attribute of an *Employees* relation. Suppose that a certain employee's mobile phone number is missing. How do you represent this fact in the database? In a three-valued logic implementation, the mobile phone attribute should allow the use of a special marker for a missing value. Then a predicate comparing the mobile phone attribute with some specific number will yield *unknown* for the case with the missing value. Three-valued predicate logic refers to the three possible logical values that can result from a predicate—*true*, *false*, and *unknown*.

Some people believe that three-valued predicate logic is nonrelational, whereas others believe that it is relational. Codd actually advocated for four-valued predicate logic, saying that there were two different cases of missing values: missing but applicable (A-Values marker), and missing but inapplicable (I-Values marker). An example of "missing but applicable" is when an employee has a mobile phone, but you don't know what the mobile phone number is. An example of "missing but inapplicable" is when an employee doesn't have a mobile phone at all. According to Codd, two special markers should be used to support these two cases of missing values. SQL implements three-valued predicate logic by supporting the *NULL* marker to signify the generic concept of a missing value. Support for *NULLs* and three-valued predicate logic in SQL is the source of a great deal of confusion and complexity, though one can argue that missing values are part of reality. In addition, the alternative—using only two-valued predicate logic—is no less problematic.

Note As mentioned, a *NULL* is not a value but rather a marker for a missing value. Therefore, though unfortunately it's common, the use of the terminology "*NULL* value" is incorrect. The correct terminology is "*NULL* marker" or just "*NULL*." In the book, I chose to use the latter because it's more common in the SQL community.

Constraints

One of the greatest benefits of the relational model is the ability to define data integrity as part of the model. Data integrity is achieved through rules called *constraints* that are defined in the data model and enforced by the RDBMS. The simplest methods of enforcing integrity are assigning an attribute type with its attendant "nullability" (whether it supports or doesn't support *NULLs*). Constraints are also enforced through the model itself; for example, the relation *Orders(orderid, orderdate, duedate, shipdate)* allows three distinct dates per order, whereas the relations *Employees(empid)* and *EmployeeChildren(empid, childname)* allow zero to countable infinity children per employee.

Other examples of constraints include *candidate keys,* which provide entity integrity, and *foreign keys,* which provide referential integrity. A candidate key is a key defined on one or more attributes that prevents more than one occurrence of the same tuple (*row* in SQL) in a relation. A predicate based on a candidate key can uniquely identify a row (such as an employee). You can define multiple candidate keys in a relation. For example, in an *Employees* relation, you can define candidate keys on *employeeid*, on *SSN* (Social Security number), and others. Typically, you arbitrarily choose one of the candidate keys as the *primary key* (for example, *employeeid* in the *Employees* relation) and use that as the preferred way to identify a row. All other candidate keys are known as *alternate keys*.

Foreign keys are used to enforce referential integrity. A foreign key is defined on one or more attributes of a relation (known as the *referencing relation*) and references a candidate key in another (or possibly the same) relation. This constraint restricts the values in the referencing relation's foreign-key attributes to the values that appear in the referenced relation's candidate-key attributes. For example, suppose that the *Employees* relation has a foreign key defined on the attribute *departmentid*, which references the primary-key attribute *departmentid* in the *Departments* relation. This means that the values in *Employees.departmentid* are restricted to the values that appear in *Departments*. *departmentid*.

Normalization

The relational model also defines *normalization rules* (also known as *normal forms*). Normalization is a formal mathematical process to guarantee that each entity will be represented by a single relation. In a normalized database, you avoid anomalies during data modification and keep redundancy to a minimum without sacrificing completeness. If you follow Entity Relationship Modeling (ERM), and represent each entity and its attributes, you probably won't need normalization; instead, you will apply normalization only to reinforce and ensure that the model is correct. You can find the definition of ERM in the following Wikipedia article: *https://en.wikipedia.org/wiki/Entity%E2%80%93relationship_model*.

The following sections briefly cover the first three normal forms (1NF, 2NF, and 3NF) introduced by Codd.

1NF The first normal form says that the tuples (rows) in the relation (table) must be unique and attributes should be atomic. This is a redundant definition of a relation; in other words, if a table truly represents a relation, it is already in first normal form.

You achieve unique rows in SQL by defining a unique key for the table.

You can operate on attributes only with operations that are defined as part of the attribute's type. Atomicity of attributes is subjective in the same way that the definition of a set is subjective. As an example, should an employee name in an *Employees* relation be expressed with one attribute (*fullname*), two (*firstname* and *lastname*), or three (*firstname*, *middlename*, and *lastname*)? The answer depends on the application. If the application needs to manipulate the parts of the employee's name separately (such as for search purposes), it makes sense to break them apart; otherwise, it doesn't.

In the same way that an attribute might not be atomic enough based on the needs of the applications that use it, an attribute might also be subatomic. For example, if an address attribute is considered atomic for the applications that use it, not including the city as part of the address would violate the first normal form.

This normal form is often misunderstood. Some people think that an attempt to mimic arrays violates the first normal form. An example would be defining a *YearlySales* relation with the following attributes: *salesperson*, *qty2014*, *qty2015*, and *qty2016*. However, in this example, you don't really violate the first normal form; you simply impose a constraint—restricting the data to three specific years: 2014, 2015, and 2016.

2NF The second normal form involves two rules. One rule is that the data must meet the first normal form. The other rule addresses the relationship between nonkey and candidate-key attributes. For every candidate key, every nonkey attribute has to be fully functionally dependent on the entire candidate key. In other words, a nonkey attribute cannot be fully functionally dependent on part of a candidate key. To put it more informally, if you need to obtain any nonkey attribute value, you need to provide the values of all attributes of a candidate key from the same tuple. You can find any value of any attribute of any tuple if you know all the attribute values of a candidate key.

As an example of violating the second normal form, suppose that you define a relation called *Orders* that represents information about orders and order lines. (See Figure 1-1.) The *Orders* relation contains the following attributes: *orderid*, *productid*, *orderdate*, *qty*, *customerid*, and *companyname*. The primary key is defined on *orderid* and *productid*.

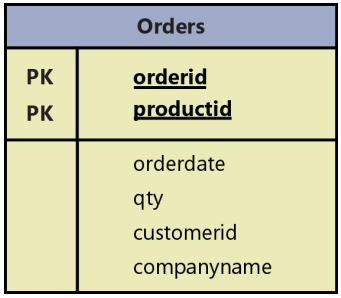

FIGURE 1-1 Data model before applying 2NF.

The second normal form is violated in Figure 1-1 because there are nonkey attributes that depend on only part of a candidate key (the primary key, in this example). For example, you can find the *orderdate* of an order, as well as *customerid* and *companyname*, based on the *orderid* alone.

To conform to the second normal form, you would need to split your original relation into two relations: *Orders* and *OrderDetails* (as shown in Figure 1-2). The *Orders* relation would include the attributes *orderid*, *orderdate*, *customerid*, and *companyname*, with the primary key defined on *orderid*. The *OrderDetails* relation would include the attributes *orderid*, *productid*, and *qty*, with the primary key defined on *orderid* and *productid*.

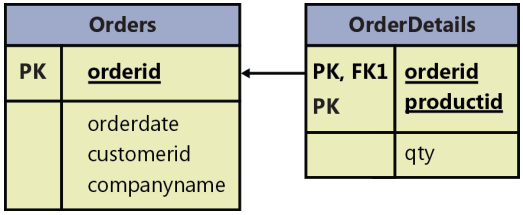

FIGURE 1-2 Data model after applying 2NF and before 3NF.

3NF The third normal form also has two rules. The data must meet the second normal form. Also, all nonkey attributes must be dependent on candidate keys nontransitively. Informally, this rule means that all nonkey attributes must be mutually independent. In other words, one nonkey attribute cannot be dependent on another nonkey attribute.

The *Orders* and *OrderDetails* relations described previously now conform to the second normal form. Remember that the *Orders* relation at this point contains the attributes *orderid*, *orderdate*, *customerid*, and *companyname*, with the primary key defined on *orderid*. Both *customerid* and *companyname* depend on the whole primary key—*orderid*. For example, you need the entire primary key to find the *customerid* representing the customer who placed the order. Similarly, you need the whole primary key to find the company name of the customer who placed the order. However, *customerid* and *companyname* are also dependent on each other. To meet the third normal form, you need to add a *Customers* relation (shown in Figure 1-3) with the attributes *customerid* (as the primary key) and *companyname*. Then you can remove the *companyname* attribute from the *Orders* relation.

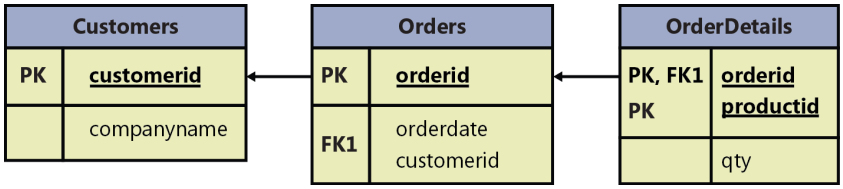

FIGURE 1-3 Data model after applying 3NF.

Informally, 2NF and 3NF are commonly summarized with the sentence, "Every non-key attribute is dependent on the key, the whole key, and nothing but the key—so help me Codd."

There are higher normal forms beyond Codd's original first three normal forms that involve compound primary keys and temporal databases, but they are outside the scope of this book.

Note SQL, as well as T-SQL, permit violating all the normal forms in real tables. It's the data modeler's prerogative and responsibility to design a normalized model.

Types of database systems

Two main types of systems, or workloads, use SQL Server as their database and T-SQL to manage and manipulate the data: online transactional processing (OLTP) and data warehouses (DWs). Figure 1-4 illustrates those systems and the transformation process that usually takes place between them.

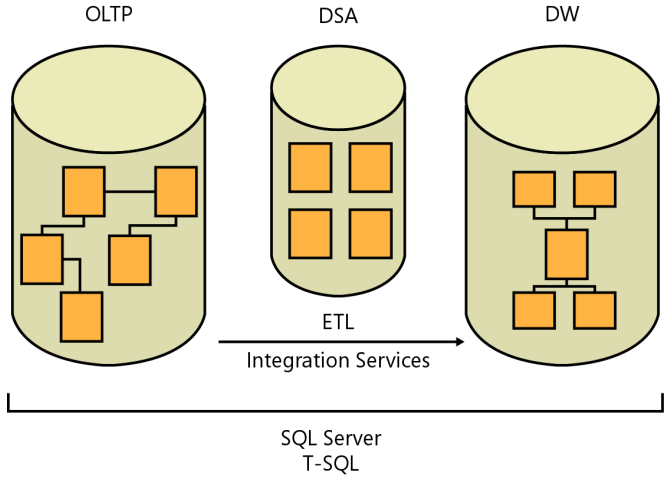

FIGURE 1-4 Classes of database systems.

Here's a quick description of what each acronym represents:

- OLTP: online transactional processing
- DSA: data-staging area
- DW: data warehouse
- ETL: extract, transform, and load

Online transactional processing

Data is entered initially into an online transactional processing system. The primary focus of an OLTP system is data entry and not reporting—transactions mainly insert, update, and delete data. The relational model is targeted primarily at OLTP systems, where a normalized model provides both good performance for data entry and data consistency. In a normalized environment, each table represents a single entity and keeps redundancy to a minimum. When you need to modify a fact, you need to modify it in only one place. This results in optimized performance for data modifications and little chance for error.

However, an OLTP environment is not suitable for reporting purposes because a normalized model usually involves many tables (one for each entity) with complex relationships. Even simple reports require joining many tables, resulting in complex and poorly performing queries.

You can implement an OLTP database in SQL Server and both manage it and query it with T-SQL.

Data warehouses

A *data warehouse* (DW) is an environment designed for data-retrieval and reporting purposes. When it serves an entire organization, such an environment is called a *data warehouse*; when it serves only part of the organization (such as a specific department) or a subject matter area in the organization, it is called a *data mart*. The data model of a data warehouse is designed and optimized mainly to support data-retrieval needs. The model has intentional redundancy, fewer tables, and simpler relationships, ultimately resulting in simpler and more efficient queries than an OLTP environment.

The simplest data-warehouse design is called a *star schema*. The star schema includes several dimension tables and a fact table. Each dimension table represents a subject by which you want to analyze the data. For example, in a system that deals with orders and sales, you will probably want to analyze data by dimensions such as customers, products, employees, and time.

In a star schema, each dimension is implemented as a single table with redundant data. For example, a product dimension could be implemented as a single *ProductDim* table instead of three normalized tables: *Products*, *ProductSubCategories*, and *ProductCategories*. If you normalize a dimension table, which results in multiple tables representing that dimension, you get what's known as a *snowflake dimension*. A schema that contains snowflake dimensions is known as a *snowflake schema*. A star schema is considered a special case of a snowflake schema.

The fact table holds the facts and measures, such as quantity and value, for each relevant combination of dimension keys. For example, for each relevant combination of customer, product,

employee, and day, the fact table would have a row containing the quantity and value. Note that data in a data warehouse is typically preaggregated to a certain level of granularity (such as a day), unlike data in an OLTP environment, which is usually recorded at the transaction level.

Historically, early versions of SQL Server mainly targeted OLTP environments, but eventually SQL Server also started targeting data-warehouse systems and data-analysis needs. You can implement a data warehouse as a SQL Server database and manage and query it with T-SQL.

The process that pulls data from source systems (OLTP and others), manipulates it, and loads it into the data warehouse is called *extract, transform, and load*, or *ETL*. SQL Server provides a tool called Microsoft SQL Server Integration Services (SSIS) to handle ETL needs.

Often the ETL process will involve the use of a data-staging area (DSA) between the OLTP and the DW. The DSA usually resides in a relational database, such as a SQL Server database, and is used as the data-cleansing area. The DSA is not open to end users.

SQL Server architecture

This section will introduce you to the SQL Server architecture, the different RDBMS flavors that Microsoft offers, the entities involved—SQL Server instances, databases, schemas, and database objects—and the purpose of each entity.

The ABCs of Microsoft RDBMS flavors

Initially, Microsoft offered mainly one enterprise-level RDBMS—an on-premises flavor called Microsoft SQL Server. These days, Microsoft offers an overwhelming plethora of options as part of its data platform, which constantly keeps evolving. Within its data platform, Microsoft offers three main RDBMS flavors, which you can think of as the *ABC flavors*: A for Appliance, B for Box (on-premises), and C for Cloud.

Box

The box, or on-premises RDBMS flavor, that Microsoft offers is called Microsoft SQL Server, or just SQL Server. This is the traditional flavor, usually installed on the customer's premises. The customer is responsible for everything—getting the hardware, installing the software, patching, high availability and disaster recovery, security, and everything else.

The customer can install multiple instances of the product on the same server (more on this in the next section) and can write queries that interact with multiple databases. It is also possible to switch the connection between databases, unless one of them is a contained database (defined later).

The querying language used is T-SQL. You can run all the code samples and exercises in this book on an on-premises SQL Server implementation, if you want. See the Appendix for details about obtaining and installing an evaluation edition of SQL Server, as well as creating the sample database.

Appliance

The idea behind the appliance flavor is to provide the customer a complete turnkey solution with preconfigured hardware and software. Speed is achieved by things being co-located, with the storage being close to the processing. The appliance is hosted locally at the customer site. Microsoft partners with hardware vendors such as Dell and HP to provide the appliance offering. Experts from Microsoft and the hardware vendor handle the performance, security, and availability aspects for the customer.

There are several appliances available today, one of which is the Microsoft Analytics Platform System (APS), which focuses on data warehousing and big data technologies. This appliance includes a data-warehouse engine called Parallel Data Warehouse (PDW), which implements massively parallel processing (MPP) technology. It also includes HDInsight, which is Microsoft's Hadoop distribution (big data solution). APS also includes a querying technology called PolyBase, which allows using T-SQL queries across relational data from PDW and nonrelational data from HDInsight.

Cloud

Cloud computing provides computing resources on demand from a shared pool of resources. Microsoft's RDBMS technologies can be provided both as private-cloud and public-cloud services. A *private cloud* is cloud infrastructure that services a single organization and usually uses virtualization technology. It's typically hosted locally at the customer site, and maintained by the IT group in the organization. It's about self-service agility, allowing the users to deploy resources on demand. It provides standardization and usage metering. The database engine is usually a box engine (hence the same T-SQL is used to manage and manipulate the data).

As for the public cloud, the services are provided over the network and available to the public. Microsoft provides two forms of public RDBMS cloud services: infrastructure as a service (IaaS) and platform as a service (PaaS). With IaaS, you provision a virtual machine (VM) that resides in Microsoft's cloud infrastructure. As a starting point, you can choose between several preconfigured VMs that already have a certain version and edition of SQL Server (box engine) installed on them. The hardware is maintained by Microsoft, but you're responsible for maintaining and patching the software. It's essentially like maintaining your own SQL Server installation—one that happens to reside on Microsoft's hardware.

With PaaS, Microsoft provides the database cloud platform as a service. It's hosted in Microsoft's data centers. Hardware, software installation and maintenance, high availability and disaster recovery, and patching are all responsibilities of Microsoft. The customer is still responsible for index and query tuning, however.

Microsoft provides a number of PaaS database offerings. For OLTP systems, it offers the Azure SQL Database service. It's also referred to more shortly as just SQL Database. The customer can have multiple databases on the cloud server (a conceptual server, of course) but cannot switch between databases.

Interestingly, Microsoft uses the same code base for SQL Database and SQL Server. So most of the T-SQL language surface is exposed (eventually) in both environments in the same manner. Therefore, most of the T-SQL you'll learn about in this book is applicable to both environments. You can read about the differences that do exist here: *https://azure.microsoft.com/en-us/documentation/articles/sql-database-transact-sql-*

information. You should also note that the update and deployment rate of new versions of SQL Database are faster than that of the on-premises SQL Server. Therefore, some T-SQL features might be exposed in SQL Database before they show up in the on-premises SQL Server version.

Microsoft also provides a PaaS offering for data-warehouse systems called Microsoft Azure SQL Data Warehouse (also called Azure SQL Data Warehouse or just SQL Data Warehouse). This service is basically PDW/APS in the cloud. Microsoft uses the same code base for both the appliance and the cloud service. You manage and manipulate data in APS and SQL Data Warehouse with T-SQL, although it's not the same T-SQL surface as in SQL Server and SQL Database, yet.

Microsoft also offers other cloud data services, such as Data Lake for big data–related services, Azure DocumentDB for NoSQL document database services, and others.

Confused? If it's any consolation, you're not alone. Like I said, Microsoft provides an overwhelming plethora of database-related technologies. Curiously, the one thread that is common to many of them is T-SQL.

SQL Server instances

In the box product, an instance of SQL Server, as illustrated in Figure 1-5, is an installation of a SQL Server database engine or service. You can install multiple instances of on-premises SQL Server on the same computer. Each instance is completely independent of the others in terms of security and the data that it manages, and in all other respects. At the logical level, two different instances residing on the same computer have no more in common than two instances residing on two separate computers. Of course, same-computer instances do share the server's physical resources, such as CPU, memory, and disk.

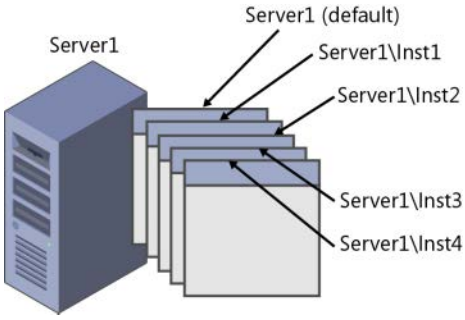

FIGURE 1-5 Multiple instances of SQL Server on the same computer.

You can set up one of the multiple instances on a computer as the *default instance*, whereas all others must be *named instances*. You determine whether an instance is the default or a named one upon installation; you cannot change that decision later. To connect to a default instance, a client application needs to specify the computer's name or IP address. To connect to a named instance, the client needs to specify the computer's name or IP address, followed by a backslash (\), followed by the instance name (as provided upon installation). For example, suppose you have two instances of SQL Server installed on a computer called *Server1*. One of these instances was installed as the default

instance, and the other was installed as a named instance called *Inst1*. To connect to the default instance, you need to specify only *Server1* as the server name. However, to connect to the named instance, you need to specify both the server and the instance name: *Server1\Inst1*.

There are various reasons why you might want to install multiple instances of SQL Server on the same computer, but I'll mention only a couple here. One reason is to save on support costs. For example, to test the functionality of features in response to support calls or reproduce errors that users encounter in the production environment, the support department needs local installations of SQL Server that mimic the user's production environment in terms of version, edition, and service pack of SQL Server. If an organization has multiple user environments, the support department needs multiple installations of SQL Server. Rather than having multiple computers, each hosting a different installation of SQL Server that must be supported separately, the support department can have one computer with multiple installed instances. Of course, you can achieve a similar result by using multiple virtual machines.

As another example, consider people like me who teach and lecture about SQL Server. For us, it is convenient to be able to install multiple instances of SQL Server on the same laptop. This way, we can perform demonstrations against different versions of the product, showing differences in behavior between versions, and so on.

As a final example, providers of database services sometimes need to guarantee their customers complete security separation of their data from other customers' data. At least in the past, the database provider could have a very powerful data center hosting multiple instances of SQL Server, rather than needing to maintain multiple less-powerful computers, each hosting a different instance. More recently, cloud solutions and advanced virtualization technologies make it possible to achieve similar goals.

Databases

You can think of a database as a container of objects such as tables, views, stored procedures, and other objects. Each instance of SQL Server can contain multiple databases, as illustrated in Figure 1-6. When you install an on-premises flavor of SQL Server, the setup program creates several system databases that hold system data and serve internal purposes. After the installation of SQL Server, you can create your own user databases that will hold application data.

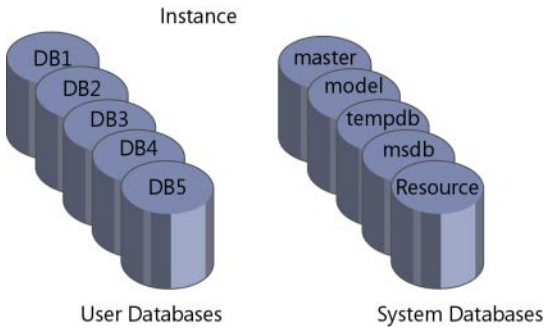

FIGURE 1-6 An example of multiple databases on a SQL Server instance.

The system databases that the setup program creates include *master*, *Resource*, *model*, *tempdb*, and *msdb*. A description of each follows:

- ***master*** The *master* database holds instance-wide metadata information, the server configuration, information about all databases in the instance, and initialization information.
- ***Resource*** The *Resource* database is a hidden, read-only database that holds the definitions of all system objects. When you query system objects in a database, they appear to reside in the *sys schema* of the local database, but in actuality their definitions reside in the *Resource* database.
- ***model*** The *model* database is used as a template for new databases. Every new database you create is initially created as a copy of *model*. So if you want certain objects (such as data types) to appear in all new databases you create, or certain database properties to be configured in a certain way in all new databases, you need to create those objects and configure those properties in the *model* database. Note that changes you apply to the *model* database will not affect existing databases—only new databases you create in the future.
- ***tempdb*** The *tempdb* database is where SQL Server stores temporary data such as work tables, sort and hash table data, row versioning information, and so on. With SQL Server, you can create temporary tables for your own use, and the physical location of those temporary tables is *tempdb*. Note that this database is destroyed and re-created as a copy of the *model* database every time you restart the instance of SQL Server.
- ***msdb*** The *msdb* database is used mainly by a service called SQL Server Agent to store its data. SQL Server Agent is in charge of automation, which includes entities such as jobs, schedules, and alerts. SQL Server Agent is also the service in charge of replication. The *msdb* database also holds information related to other SQL Server features, such as Database Mail, Service Broker, backups, and more.

In an on-premises installation of SQL Server, you can connect directly to the system databases *master*, *model*, *tempdb*, and *msdb*. In SQL Database, you can connect directly only to the system database *master*. If you create temporary tables or declare table variables (more on this topic in Chapter 11, "Programmable objects"), they are created in *tempdb*, but you cannot connect directly to *tempdb* and explicitly create user objects there.

You can create multiple user databases (up to 32,767) within an instance. A user database holds objects and data for an application.

You can define a property called *collation* at the database level that will determine default language support, case sensitivity, and sort order for character data in that database. If you do not specify a collation for the database when you create it, the new database will use the default collation of the instance (chosen upon installation).

To run T-SQL code against a database, a client application needs to connect to a SQL Server instance and be in the context of, or use, the relevant database. The application can still access objects from other databases by adding the database name as a prefix.

In terms of security, to be able to connect to a SQL Server instance, the database administrator (DBA) must create a *login* for you. The login can be tied to your Microsoft Windows credentials, in which case it is called a *Windows authenticated login*. With a Windows authenticated login, you can't provide login and password information when connecting to SQL Server because you already provided those when you logged on to Windows. The login can be independent of your Windows credentials, in which case it's called a *SQL Server authenticated login*. When connecting to SQL Server using a SQL Server authenticated login, you will need to provide both a login name and a password.

The DBA needs to map your login to a *database user* in each database you are supposed to have access to. The database user is the entity that is granted permissions to objects in the database.

SQL Server supports a feature called *contained databases* that breaks the connection between a database user and an instance-level login. The user (Windows or SQL authenticated) is fully contained within the specific database and is not tied to a login at the instance level. When connecting to SQL Server, the user needs to specify the database he or she is connecting to, and the user cannot subsequently switch to other user databases.

So far, I've mainly mentioned the logical aspects of databases. If you're using SQL Database, your only concern is that logical layer. You do not deal with the physical layout of the database data and log files, *tempdb*, and so on. But if you're using a box version of SQL Server, you are responsible for the physical layer as well. Figure 1-7 shows a diagram of the physical database layout.

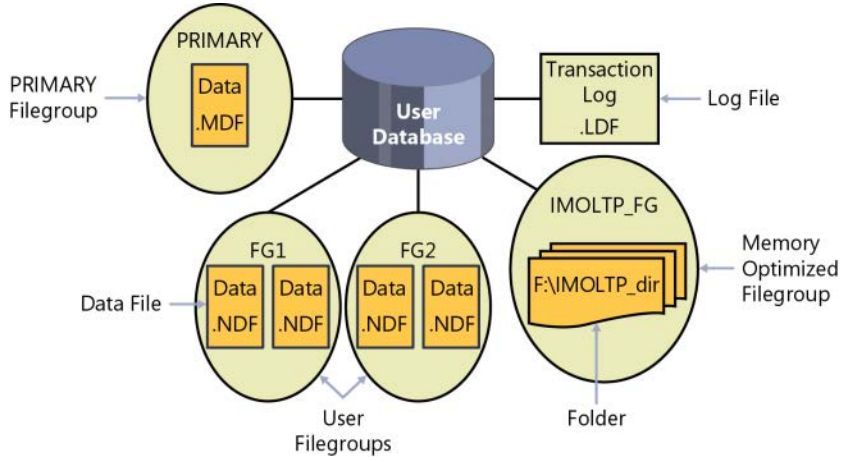

FIGURE 1-7 Database layout.

The database is made up of data files, transaction log files, and optionally checkpoint files holding memory-optimized data (part of a feature called *In-Memory OLTP*, which I describe shortly). When you create a database, you can define various properties for data and log files, including the file name, location, initial size, maximum size, and an autogrowth increment. Each database must have at least one data file and at least one log file (the default in SQL Server). The data files hold object data, and the log files hold information that SQL Server needs to maintain transactions.

Although SQL Server can write to multiple data files in parallel, it can write to only one log file at a time, in a sequential manner. Therefore, unlike with data files, having multiple log files does not result in a performance benefit. You might need to add log files if the disk drive where the log resides runs out of space.

Data files are organized in logical groups called *filegroups*. A filegroup is the target for creating an object, such as a table or an index. The object data will be spread across the files that belong to the target filegroup. Filegroups are your way of controlling the physical locations of your objects. A database must have at least one filegroup called *PRIMARY*, and it can optionally have other user filegroups as well. The *PRIMARY* filegroup contains the primary data file (which has an .mdf extension) for the database, and the database's system catalog. You can optionally add secondary data files (which have an .ndf extension) to *PRIMARY*. User filegroups contain only secondary data files. You can decide which filegroup is marked as the default filegroup. Objects are created in the default filegroup when the object creation statement does not explicitly specify a different target filegroup.

File extensions .mdf, .ldf, and .ndf

The database file extensions .mdf and .ldf are quite straightforward. The extension *.mdf* stands for *Master Data File* (not to be confused with the *master* database), and *.ldf* stands for *Log Data File*. According to one anecdote, when discussing the extension for the secondary data files, one of the developers suggested, humorously, using .ndf to represent "Not Master Data File," and the idea was accepted.

The SQL Server database engine includes a memory-optimized engine called In-Memory OLTP. You can use this feature to integrate memory-optimized objects, such as memory-optimized tables and natively compiled procedures, into your database. To do so, you need to create a filegroup in the database marked as containing memory-optimized data and, within it, at least one path to a folder. SQL Server stores checkpoint files with memory-optimized data in that folder, and it uses those to recover the data every time SQL Server is restarted.

Schemas and objects

When I said earlier that a database is a container of objects, I simplified things a bit. As illustrated in Figure 1-8, a database contains schemas, and schemas contain objects. You can think of a schema as a container of objects, such as tables, views, stored procedures, and others.

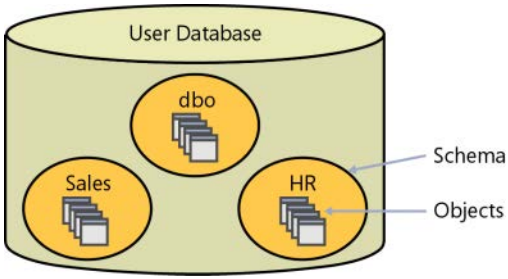

FIGURE 1-8 A database, schemas, and database objects.

You can control permissions at the schema level. For example, you can grant a user SELECT permissions on a schema, allowing the user to query data from all objects in that schema. So security is one of the considerations for determining how to arrange objects in schemas.

The schema is also a namespace—it is used as a prefix to the object name. For example, suppose you have a table named *Orders* in a schema named *Sales*. The schema-qualified object name (also known as the *two-part object name*) is *Sales.Orders*. You can refer to objects in other databases by adding the database name as a prefix (*three-part object name*), and to objects in other instances by adding the instance name as a prefix (*four-part object name*). If you omit the schema name when referring to an object, SQL Server will apply a process to resolve the schema name, such as checking whether the object exists in the user's default schema and, if the object doesn't exist, checking whether it exists in the *dbo* schema. Microsoft recommends that when you refer to objects in your code you always use the two-part object names. There are some relatively insignificant extra costs involved in resolving the schema name when you don't specify it explicitly. But as insignificant as this extra cost might be, why pay it? Also, if multiple objects with the same name exist in different schemas, you might end up getting a different object than the one you wanted.

Creating tables and defining data integrity

This section describes the fundamentals of creating tables and defining data integrity using T-SQL. Feel free to run the included code samples in your environment.

More Info If you don't know yet how to run code against SQL Server, the Appendix will help you get started.

As mentioned earlier, DML rather than DDL is the focus of this book. Still, you need to understand how to create tables and define data integrity. I won't go into the explicit details here, but I'll provide a brief description of the essentials.

Before you look at the code for creating a table, remember that tables reside within schemas, and schemas reside within databases. The examples use the book's sample database, *TSQLV4*, and a schema called *dbo*.

More Info See the Appendix for details on creating the sample database.

The examples here use a schema named *dbo* that is created automatically in every database and is also used as the default schema for users who are not explicitly associated with a different schema.

Creating tables

The following code creates a table named *Employees* in the *dbo* schema in the *TSQLV4* database:

```
USE TSQLV4;

DROP TABLE IF EXISTS dbo.Employees;

CREATE TABLE dbo.Employees
(
  empid     INT         NOT NULL,
  firstname VARCHAR(30) NOT NULL,
  lastname  VARCHAR(30) NOT NULL,
  hiredate  DATE        NOT NULL,
  mgrid     INT         NULL,
  ssn       VARCHAR(20) NOT NULL,
  salary    MONEY       NOT NULL
);
```

The *USE* statement sets the current database context to that of *TSQLV4*. It is important to incorporate the *USE* statement in scripts that create objects to ensure that SQL Server creates the objects in the specified database. In an on-premises SQL Server implementation, the *USE* statement can actually change the database context from one to another. In SQL Database, you cannot switch between different databases, but the *USE* statement will not fail as long as you are already connected to the target database. So even in SQL Database, I recommend having the *USE* statement to ensure that you are connected to the right database when creating your objects.

The *DROP IF EXISTS* command drops the table if it already exists. Note that this command was introduced in SQL Server 2016. If you're using earlier versions of SQL Server, use the following statement instead:

```
IF OBJECT_ID(N'dbo.Employees', N'U') IS NOT NULL DROP TABLE dbo.Employees;
```

The *IF* statement invokes the *OBJECT_ID* function to check whether the *Employees* table already exists in the current database. The *OBJECT_ID* function accepts an object name and type as inputs. The type *U* represents a user table. This function returns the internal object ID if an object with the specified input name and type exists, and *NULL* otherwise. If the function returns a *NULL*, you know that the object doesn't exist. In our case, the code drops the table if it already exists and then creates a new one. Of course, you can choose a different treatment, such as simply not creating the object if it already exists.

The *CREATE TABLE* statement is in charge of defining what I referred to earlier as the heading of the relation. Here you specify the name of the table and, in parentheses, the definition of its attributes (columns).

Notice the use of the two-part name *dbo.Employees* for the table name, as recommended earlier. If you omit the schema name, for ad-hoc queries SQL Server will assume the default schema associated with the database user running the code. For queries in stored procedures, SQL Server will assume the schema associated with the procedure's owner.

For each attribute, you specify the attribute name, data type, and whether the value can be *NULL* (which is called *nullability*).

In the *Employees* table, the attributes *empid* (employee ID) and *mgrid* (manager ID) are each defined with the *INT* (four-byte integer) data type; the *firstname*, *lastname*, and *ssn* (US Social Security number) are defined as *VARCHAR* (variable-length character string with the specified maximum supported number of characters); and *hiredate* is defined as *DATE* and *salary* is defined as *MONEY*.

If you don't explicitly specify whether a column allows or disallows *NULLs*, SQL Server will have to rely on defaults. Standard SQL dictates that when a column's nullability is not specified, the assumption should be *NULL* (allowing *NULLs*), but SQL Server has settings that can change that behavior. I recommend that you be explicit and not rely on defaults. Also, I recommend defining a column as *NOT NULL* unless you have a compelling reason to support *NULLs*. If a column is not supposed to allow *NULLs* and you don't enforce this with a *NOT NULL* constraint, you can rest assured that *NULLs* will occur. In the *Employees* table, all columns are defined as *NOT NULL* except for the *mgrid* column. A *NULL* in the *mgrid* column would represent the fact that the employee has no manager, as in the case of the CEO of the organization.

Coding style

You should be aware of a few general notes regarding coding style, the use of white spaces (space, tab, new line, and so on), and semicolons. I'm not aware of any formal coding styles. My advice is that you use a style that you and your fellow developers feel comfortable with. What ultimately matters most is the consistency, readability, and maintainability of your code. I have tried to reflect these aspects in my code throughout the book.

T-SQL lets you use white spaces quite freely in your code. You can take advantage of white space to facilitate readability. For example, I could have written the code in the previous section as a single line. However, the code wouldn't have been as readable as when it is broken into multiple lines that use indentation.

The practice of using a semicolon to terminate statements is standard and, in fact, is a requirement in several other database platforms. SQL Server requires the semicolon only in particular cases—but in cases where a semicolon is not required, using one doesn't cause problems. I recommend that you adopt the practice of terminating all statements with a semicolon. Not only will doing this improve the readability of your code, but in some cases it can save you some grief. (When a semicolon is required and is *not* specified, the error message SQL Server produces is not always clear.)

Note The SQL Server documentation indicates that not terminating T-SQL statements with a semicolon is a deprecated feature. This means that the long-term goal is to enforce use of the semicolon in a future version of the product. That's one more reason to get into the habit of terminating all your statements, even where it's currently not required.

Defining data integrity

As mentioned earlier, one of the great benefits of the relational model is that data integrity is an integral part of it. Data integrity enforced as part of the model—namely, as part of the table definitions—is considered *declarative data integrity*. Data integrity enforced with code—such as with stored procedures or triggers—is considered *procedural data integrity*.

Data type and nullability choices for attributes and even the data model itself are examples of declarative data integrity constraints. In this section, I will describe other examples of declarative constraints: primary key, unique, foreign key, check, and default constraints. You can define such constraints when creating a table as part of the *CREATE TABLE* statement, or you can define them for already-created tables by using an *ALTER TABLE* statement. All types of constraints except for default constraints can be defined as *composite constraints*—that is, based on more than one attribute.

Primary-key constraints

A primary-key constraint enforces the uniqueness of rows and also disallows *NULLs* in the constraint attributes. Each unique set of values in the constraint attributes can appear only once in the table—in other words, only in one row. An attempt to define a primary-key constraint on a column that allows *NULLs* will be rejected by the RDBMS. Each table can have only one primary key.

Here's an example of defining a primary-key constraint on the *empid* attribute in the *Employees* table that you created earlier:

```
ALTER TABLE dbo.Employees
  ADD CONSTRAINT PK_Employees
  PRIMARY KEY(empid);
```

With this primary key in place, you can be assured that all *empid* values will be unique and known. An attempt to insert or update a row such that the constraint would be violated will be rejected by the RDBMS and result in an error.

To enforce the uniqueness of the logical primary-key constraint, SQL Server will create a unique index behind the scenes. A *unique index* is a physical mechanism used by SQL Server to enforce uniqueness. Indexes (not necessarily unique ones) are also used to speed up queries by avoiding unnecessary full table scans (similar to indexes in books).

Unique constraints

A unique constraint enforces the uniqueness of rows, allowing you to implement the concept of alternate keys from the relational model in your database. Unlike with primary keys, you can define multiple unique constraints within the same table. Also, a unique constraint is not restricted to columns defined as *NOT NULL*.

The following code defines a unique constraint on the *ssn* column in the *Employees* table:

```
ALTER TABLE dbo.Employees
  ADD CONSTRAINT UNQ_Employees_ssn
  UNIQUE(ssn);
```

As with a primary-key constraint, SQL Server will create a unique index behind the scenes as the physical mechanism to enforce the logical unique constraint.

According to standard SQL, a column with a unique constraint is supposed to allow multiple *NULLs* (as if two *NULLs* were different from each other). However, SQL Server's implementation rejects duplicate *NULLs* (as if two *NULLs* were equal to each other). To emulate the standard unique constraint in SQL Server you can use a unique filtered index that filters only non-*NULL* values. For example, suppose that the column *ssn* allowed *NULLs,* and you wanted to create such an index instead of a unique constraint. You would have used the following code:

```
CREATE UNIQUE INDEX idx_ssn_notnull ON dbo.Employees(ssn) WHERE ssn IS NOT NULL;
```

The index is defined as a unique one, and the filter excludes *NULLs* from the index, so duplicate *NULLs* will be allowed, whereas duplicate non-*NULL* values won't be allowed.

Foreign-key constraints

A foreign-key enforces referential integrity. This constraint is defined on one or more attributes in what's called the *referencing* table and points to candidate-key (primary-key or unique-constraint) attributes in what's called the *referenced* table. Note that the referencing and referenced tables can be one and the same. The foreign key's purpose is to restrict the values allowed in the foreign-key columns to those that exist in the referenced columns.

The following code creates a table called *Orders* with a primary key defined on the *orderid* column:

```
DROP TABLE IF EXISTS dbo.Orders;

CREATE TABLE dbo.Orders
(
  orderid   INT         NOT NULL,
  empid     INT         NOT NULL,
  custid    VARCHAR(10) NOT NULL,
  orderts   DATETIME2   NOT NULL,
  qty       INT         NOT NULL,
  CONSTRAINT PK_Orders
    PRIMARY KEY(orderid)
);
```

Suppose you want to enforce an integrity rule that restricts the values supported by the *empid* column in the *Orders* table to the values that exist in the *empid* column in the *Employees* table. You can achieve this by defining a foreign-key constraint on the *empid* column in the *Orders* table pointing to the *empid* column in the *Employees* table, like the following:

```
ALTER TABLE dbo.Orders
  ADD CONSTRAINT FK_Orders_Employees
  FOREIGN KEY(empid)
  REFERENCES dbo.Employees(empid);
```

Similarly, if you want to restrict the values supported by the *mgrid* column in the *Employees* table to the values that exist in the *empid* column of the same table, you can do so by adding the following foreign key:

```
ALTER TABLE dbo.Employees
  ADD CONSTRAINT FK_Employees_Employees
  FOREIGN KEY(mgrid)
  REFERENCES dbo.Employees(empid);
```

Note that *NULLs* are allowed in the foreign-key columns (*mgrid* in the last example) even if there are no *NULLs* in the referenced candidate-key columns.

The preceding two examples are basic definitions of foreign keys that enforce a referential action called *no action*. No action means that attempts to delete rows from the referenced table or update the referenced candidate-key attributes will be rejected if related rows exist in the referencing table. For example, if you try to delete an employee row from the *Employees* table when there are related orders in the *Orders* table, the RDBMS will reject such an attempt and produce an error.

You can define the foreign key with actions that will compensate for such attempts (to delete rows from the referenced table or update the referenced candidate-key attributes when related rows exist in the referencing table). You can define the options *ON DELETE* and *ON UPDATE* with actions such as *CASCADE*, *SET DEFAULT*, and *SET NULL* as part of the foreign-key definition. *CASCADE* means that the operation (delete or update) will be cascaded to related rows. For example, *ON DELETE CASCADE* means that when you delete a row from the referenced table, the RDBMS will delete the related rows from the referencing table. *SET DEFAULT* and *SET NULL* mean that the compensating action will set the foreign-key attributes of the related rows to the column's default value or *NULL*, respectively. Note that regardless of which action you choose, the referencing table will have only orphaned rows in the case of the exception with *NULLs* that I mentioned earlier. Parents with no children are always allowed.

Check constraints

You can use a *check constraint* to define a predicate that a row must meet to be entered into the table or to be modified. For example, the following check constraint ensures that the salary column in the *Employees* table will support only positive values:

```
ALTER TABLE dbo.Employees
  ADD CONSTRAINT CHK_Employees_salary
  CHECK(salary > 0.00);
```

An attempt to insert or update a row with a nonpositive salary value will be rejected by the RDBMS. Note that a check constraint rejects an attempt to insert or update a row when the predicate evaluates to *FALSE*. The modification will be accepted when the predicate evaluates to either *TRUE* or *UNKNOWN*. For example, salary –1000 will be rejected, whereas salaries 50000 and *NULL* will both be accepted (if the column allowed *NULLs*). As mentioned earlier, SQL is based on three-valued logic, which results in two actual actions. With a check constraint, the row is either accepted or rejected.

When adding check and foreign-key constraints, you can specify an option called *WITH NOCHECK* that tells the RDBMS you want it to bypass constraint checking for existing data. This is considered a bad practice because you cannot be sure your data is consistent. You can also disable or enable existing check and foreign-key constraints.

Default constraints

A default constraint is associated with a particular attribute. It's an expression that is used as the default value when an explicit value is not specified for the attribute when you insert a row. For example, the following code defines a default constraint for the *orderts* attribute (representing the order's time stamp):

```
ALTER TABLE dbo.Orders
  ADD CONSTRAINT DFT_Orders_orderts
  DEFAULT(SYSDATETIME()) FOR orderts;
```

The default expression invokes the *SYSDATETIME* function, which returns the current date and time value. After this default expression is defined, whenever you insert a row in the *Orders* table and do not explicitly specify a value in the *orderts* attribute, SQL Server will set the attribute value to *SYSDATETIME*.

When you're done, run the following code for cleanup:

```
DROP TABLE IF EXISTS dbo.Orders, dbo.Employees;
```

Conclusion

This chapter provided a brief background to T-SQL querying and programming. It presented a theoretical background, explaining the strong foundations that T-SQL is based on. It gave an overview of the SQL Server architecture and concluded with sections that demonstrated how to use T-SQL to create tables and define data integrity. I hope that by now you see that there's something special about SQL, and that it's not just a language that can be learned as an afterthought. This chapter equipped you with fundamental concepts—the actual journey is just about to begin.

This page intentionally left blank

Index

Symbols and Numbers

- (negative), 51–52
- (subtraction), 51–52
% (percent) wildcard, 71–72
(), parentheses, 51–52
* (asterisk), SELECT clauses and, 41
* (multiplication), 51–52
, , = (comparison operators), 51–52
.ldf (Log Data File), 18
.mdf (Master Data File), 18
.ndf (Not Master Data File), 18
@@ identity, 255–257
_ (underscore) wildcard, 72
+ (addition), 51–52
+ (concatenation), 51–52
+ (positive), 51–52
= (assignment), 51–52
[^character list or range] wildcard, 73
[character-character] wildcard, 72–73
[list of characters] wildcard, 72
1NF, 8
2NF, 8–10
3NF, 9–10

A

ABC flavors (Appliance, Box, Cloud), 12–14
accent sensitivity, character data, 61
ACID (automaticity, consistency, isolation, and durability), 320–322
addition (+), 51–52
administrative tasks, dynamic SQL and, 379–383
aggregate functions. *See also* COUNT; also pivoting data; also window functions
  aggregate window functions, 221–224
  NULLs and, 33–35
  running aggregates, subqueries and, 144–145
aliases
  common table expressions (CTEs), assigning in, 167–168
  cross joins and, 104
  derived tables and, 162–164
  derived tables, nesting and, 165–166
  inline table-valued functions (TVFs), 179
  ORDER BY clause and, 42
  SELECT clauses and, 36–41
  self joins and, 106
ALL, duplicates and, 194
all-at-once operations, 58–59
ALTER DATABASE, quoted identifiers, 62–63
ALTER PROC, 385
ALTER SEQUENCE, 259
ALTER TABLE
  defining data integrity, 22–25
  sequence objects and, 261–262
ALTER TABLE ADD CONSTRAINT, 264–266
ALTER TABLE DROP CONSTRAINT, 264–266
ALTER VIEW, 172
American National Standards Institute (ANSI), 2
AND
  MERGE and, 275
  precedence, 51–52
  predicates and, 49
  use of, 50
ANSI (American National Standards Institute), 2
Appliance, Box, Cloud (ABC flavors), 12–14
application-time period tables, 297. See also temporal tables
APPLY
  table expressions, overview, 180–183
  unpivoting data, 231–233
arguments, derived table queries and, 164–165
arithmetic operators, 50–52
AS clause, 36–41
  derived tables and, 161–162

assignment (=), 51–52
asterisk (*), SELECT clauses and, 41
AT TIME ZONE, 84–85
  temporal table queries, 310
attributes
  filtering in outer joins, 118–119
  normalization forms, 7–10
  propositions, predicates, and relations, 5–6
auditing, DDL triggers (data definition events), 388–390. See also temporal tables
automaticity, transactions, 320–322
automating administrative tasks, dynamic SQL and, 379–383
autonumbering
  assignment UPDATE and, 270–271
AVG, 33–35. See also window functions
Azure SQL Data Warehouse, 14
Azure SQL Database
  ABC flavors (Appliance, Box, Cloud), 13–14
  global temporary tables, 376
  isolation, overview, 332–333
  locks and blocking, 322–323
  READ COMMITTED SNAPSHOT, 343
  row-versioning-based isolation, 339
  source code, downloading of, 402–404
  website link and access information, 395–396

B

bag, 3
batches
  as unit of parsing, 364
  as unit of resolution, 366
  dynamic SQL, 379–383
  overview of, 363–364
  variables and, 361, 365
  vs. transactions, 363
BEGIN, 368
BEGIN TRAN (TRANSACTION), 319–322
BETWEEN
  precedence, 51–52
  temporal table queries, 308
  use of, 49
binding, batches and, 366
bitemporal, 297. See also temporal tables
blocking. See isolation, transactions; locks
blocking chain, 328
blocking_session_id, 330–331
BULK INSERT, 253

C

CACHE, sequence objects and, 258–259
Cancel Executing Query, 331–332
candidate keys. See also constraints
  normalization forms, 8–9
  overview of, 7
Cantor, Georg, 3
CASCADE, 24
CASE expressions
  overview of, 52–54
  pivoting with grouped queries, 226–227
case sensitivity, character data, 61
CAST
  correlated subqueries and, 141
  date and time functions, 81–83
  date and time literals, 77–78
catalog views, metadata queries, 89–91
CATCH, error handling, 390–394
change management, DDL triggers (data definition events), 388–390
CHAR, 50, 60–61
character data
  collation, 61–63
  data types, overview, 60–61
  quoted identifiers, 62–63
CHARINDEX, 66
CHECK
  CASE expressions, overview of, 52–54
  working with date and time separately, 78–79
CHECK OPTION, views and, 177–178
CHOOSE, CASE expressions, 54
closed-world assumption, 5
cloud computing, ABC flavors (Appliance, Box, Cloud), 13–14
COALESCE, CASE expressions, 54
Codd, Edgar F., 4–5
code samples, resources for getting started
  Azures SQL Database, getting started, 395–396
  source code, downloading of, 402–404
  SQL Server Books Online, 410–413
  SQL Server installation, 396–401
  SQL Server Management Studio, working with, 404–410
coding style, T-SQL, 21–22
COLLATE, 61–62
collation, 16, 61–63
column names
  aliases, assigning in common table expressions (CTEs), 167–168

derived tables and, 162
identifier names, 30
subquery substitution errors, 147–149
COLUMNPROPERTY, 92
COMMIT TRAN
lock modes and compatibility, 322–324
transactions, overview of, 319–322
common table expressions (CTEs)
arguments, using, 167–168
column aliases, assigning, 167–168
defining multiple CTEs, 168–169
exercises, 183–188
exercises, solutions for, 188–192
multiple references in, 169
overview of, 167
recursive CTEs, 169–171
comparison operators, 50–52
composite joins, 110–111
exercises, 123–128
exercises, solutions for, 129–132
COMPRESS, 70
compression, character data and, 60–61
CONCAT, 63–65
concatenation (+), 51–52
concurrency, transactions overview, 319–322
conflict detection, SNAPSHOT and, 341–343
consistency, transactions, 320–322. See also isolation, transactions
CONSTRAINT, 264–266
constraints
check constraints, defining, 24–25
default constraints, defining, 25
foreign-key constraints, defining, 23–24
overview of, 7
primary-key constraints, defining, 22
unique constraints, defining, 23
contained databases, SQL Server architecture, 17
CONTINUE, 369–370
CONVERT
date and time functions, 81–83
date and time literals, 77–78
working with date and time separately, 79
correlated subqueries. See also subqueries
EXISTS and, 141–143
overview of, 139–141
COUNT. See also window functions
GROUP BY and, 33–35
outer joins, use with, 121–123
CREATE DATABASE, DDL triggers (data definition events), 388–390
CREATE DEFAULT, 365
CREATE FUNCTION, 365
CREATE PROCEDURE, 365
CREATE RULE, 365
CREATE SCHEMA, 365
CREATE SEQUENCE, 258–259
CREATE TABLE
DDL triggers (data definition events), 388–390
defining data integrity, 22–25
use of, 20
CREATE TRIGGER, 365
CREATE VIEW, 365
CROSS APPLY
JOIN and, 180
overview of, 180–183
unpivoting data, 232–233
cross-joins
CROSS APPLY and, 180
exercises, 123–128
exercises, solutions for, 129–132
overview of, 103–107
CUBE, 235–236
CURRENT_TIMESTAMP, 80–81
cursors, 42
as programmable objects, 370–373
custom sequence, assignment UPDATE and, 270–271
CYCLE, sequence objects and, 258–259

D

Darwen, Hugh, 4–5
data definition events (DDL triggers), 387–390
data definition language (DDL), batches and, 366
data integrity
constraints and, 7
defining, 22–25
Data Lake, 14
data manipulation events (DML triggers), 387–388
Data Manipulation Language (DML), 249
batches and, 366
nested DML, 285–287
data marts, overview of, 11–12
data types
character data, 60–61
N prefix, use of, 50
operands, precedence of, 50–52
data warehouses (DWs), 10–12
database system types, 10–12
database user, SQL Server logins, 17

DATABASEPROPERTYEX, 92
databases, SQL Server
  architecture overview, 15–18
  file extensions (.mdf, .ldf, .ndf), 18
  schemas and objects, 18–19
DATALENGTH, 65–66
DATE
  data types, 73–74
  literals, 74–78
  working with date and time separately, 78–79
date and time data
  aggregate window functions, 223–224
  AT TIME ZONE, 84–85
  CAST, CONVERT, and PARSE, 81–83
  current date and time functions, 80–81
  DATEADD, 85
  DATEDIFF and DATEDIFF_BIG, 86–87
  DATENAME, 88
  DATEPART, 87
  EOMONTH, 89
  filtering date ranges, 79–80
  FROMPARTS, 88–89
  ISDATE, 88
  literals, 74–78
  missing values, in outer joins, 116–118
  SWITCHOFFSET, 83
  types of, 73–74
  working with separately, 78–79
  YEAR, MONTH, and DAY, 87–88
Date, Chris, 4–5
DATEADD, 85
DATEDIFF and DATEDIFF_BIG, 86–87
DATEFORMAT, 75
DATENAME, 88
DATEPART, 87
DATETIME
  data types, 73–74
  literals, 74–78
  working with date and time separately, 78–79
DATETIME2
  data types, 73–74
  literals, 74–78
  temporal tables, creating, 298–301
DATETIMEOFFSET
  data types, 73–74
  literals, 74–78
  SWITCHOFFSET and, 83
DAY, 87–88
DBCC CHECKIDENT, 257
DDL triggers (data definition events), 387–390
DEADLOCK_PRIORITY, 345
deadlocks, 337, 345–348
deadly embrace deadlock, 347
declarative data integrity, 22–25
DECLARE
  temporary table variables, 377–378
  variables and, 361
DECOMPRESS, 70
default constraints, defining, 25
DELETE
  DML triggers (data manipulation events), 387–388
  exercises, 287–291
  exercises, solutions for, 291–295
  joins and, 265–266
  OUTPUT and, 282–283
  overview of, 262–266
  row-versioning isolation and, 339
  table expressions, modifying data with, 276–278
  views and, 172
DENSE_RANK, window function ranking, 216–219
derived tables
  arguments, use of, 164–165
  column aliases, assigning, 163–164
  exercises, 183–188
  exercises, solutions for, 188–192
  multiple references, 166
  nesting, 165–166
  overview of, 161–162
DESC, 43
dictionary sorting, character data, 61
dirty reads. See isolation, transactions
DISTINCT
  aggregate functions and, 34–35
  duplicates and, 39–41, 194
  EXCEPT (DISTINCT), 199–200
  INTERSECT (DISTINCT), 196–197
  multivalued subqueries and, 138
  ORDER BY clause and, 43
  UNION (DISTINCT), 194–196
  window functions, ranking of, 218–219
distinct, defined, 3
DML. See Data Manipulation Language (DML)
DML triggers (data manipulation events), 387–388
DROP CONSTRAINT, sequence objects and, 261–262
DROP IF EXISTS, table creation, 20
duplicates, DISTINCT and, 39–41, 194
durability, transactions, 320–322
DWs (data warehouses), 10–12
dynamic SQL, 379–383

E

ELSE
  CASE expressions, overview of, 52–54
  IF . . . ELSE flow element, 367–368
ELSE NULL, CASE expressions, 52–54
ENCRYPTION, views and, 175–176
END, 368
Entity Relationship Modeling (ERM), 7
EOMONTH, 89
equi and non-equi joins, 111–112
  exercises, 123–128
  exercises, solutions for, 129–132
error handling
  deadlocks and, 347–348
  programmable objects and, 390–394
  stored procedures and, 385
  transactions, overview of, 320–322
ERROR_LINE, 391
ERROR_MESSAGE, 391
ERROR_NUMBER, 391
ERROR_PROCEDURE, 391
ERROR_SEVERITY, 391
ERROR_STATE, 391
ESCAPE character, 73
ETL (extract, transform, and load), 12
EVENTDATA, DDL triggers (data definition events), 388–390
EXCEPT
  exercises, 204–208
  exercises, solutions for, 208–211
  precedence and, 201–202
  use of, 199–201
EXCEPT (DISTINCT), 199–200
EXCEPT ALL, 199–201
exclusive locks, 322–324
EXEC
  dynamic SQL, 379–381
  INSERT EXEC, 251–252
execution plan, SQL server, 380–381
exercises
  DELETE, 287–291
  DELETE, solutions for, 291–295
  grouping set, 241–244
  grouping set, solutions for, 245–248
  INSERT, 287–291
  INSERT, solutions for, 291–295
  JOIN, 123–128
  JOIN, solutions for, 129–132
  MERGE, 287–291
  MERGE, solutions for, 291–295
  OFFSET-FETCH, 287–291
  OFFSET-FETCH, solutions for, 291–295
  OUTPUT, 287–291
  OUTPUT, solutions for, 291–295
  pivoting and unpivoting data, 241–244
  pivoting and unpivoting data, solutions for, 245–248
  SELECT statements, 93–97
  SELECT statements, solutions for, 97–102
  subqueries, 150–154
  subqueries, solutions for, 154–159
  table expressions, 183–188, 287–291
  table expressions, solutions for, 188–192, 291–295
  temporal tables, 311–313
  temporal tables, solutions for, 313–317
  TOP, 287–291
  TOP, solutions for, 291–295
  transactions, isolation and locks, 348–359
  UPDATE, 287–291
  UPDATE, solutions for, 291–295
  windows functions, 241–244
  windows functions, solutions for, 245–248
exercises, resources for getting started
  Azures SQL Database, getting started, 395–396
  source code, downloading of, 402–404
  SQL Server Books Online, 410–413
  SQL Server installation, 396–401
  SQL Server Management Studio, working with, 404–410
EXISTS
  correlated subqueries and, 141–143
  subqueries, NULL trouble and, 147
extract, transform, and load (ETL), 12

F

FALSE
  check constraints and, 24–25
  IF . . . ELSE flow element, 367–368
  NULLs, overview of, 54–58
  subqueries, NULL trouble and, 146–147
  WHERE and, 32
  WHILE flow element, 368–370
filters. See also subqueries
  date ranges, 79–80
  deadlocks, avoiding, 347
  HAVING, overview of, 35–36
  OFFSET-FETCH, overview of, 46–47

filters. See also subqueries *continued*
predicate logic, overview of, 4
TOP filters, overview of, 44–46
WHERE, overview of, 31–32
WITH TIES and, 46
FIRST_VALUE, offset window functions, 219–221
flow elements, 367–370
IF . . . ELSE flow element, 367–368
WHILE flow element, 368–370
fn_helpcollations, 61–63
FOR SYSTEM_TIME CONTAINED IN, 308–310
FOR SYSTEM_TIME, temporal table queries, 304–310
FOR XML, views and ORDER BY clause, 173–174
foreign-key constraints. See also constraints
data integrity, defining, 23–24
overview of, 7
FORMAT, 69
FROM. See also JOIN
derived tables and, 161–162
logical query processing order, 28
overview of, 29–30
FROMPARTS, 88–89
functions
CHARINDEX, 66
COMPRESS and DECOMPRESS, 70
CONCAT, string concatenation, 63–65
date and time functions, 80–89
FORMAT, 69
LEFT and RIGHT, 65
LEN and DATALENGTH, 65–66
PATINDEX, 66
REPLACE, 66–67
REPLICATE, 67–68
RTRIM and LTRIM, 69
STRING_SPLIT, 71
STUFF, 68
SUBSTRING, 65
UPPER and LOWER, 68–69

G

generate globally unique identifiers (GUIDs), INSERT SELECT and, 251
GENERATED ALWAYS AS ROW END, 298–301
GENERATED ALWAYS AS ROW START, 298–301
GETDATE, 80–81
GETUTCDATE, 80–81
global temporary tables, 376–377
GO, 363–367
GROUP BY. See also grouping sets
CUBE, 236
derived tables, assigning aliases, 163–164
GROUPING SETS, 235–236
HAVING and, 35–36
logical query processing order, 28
overview of, 32–35
pivoting with grouped queries, 226–227
ROLLUP, 236–237
GROUPING, 235, 238–240
grouping sets. See also GROUP BY
CUBE, 236
exercises, 241–244
exercises, solutions for, 245–248
GROUPING and GROUPING_ID functions, 238–240
GROUPING SETS subclause, 235–236
overview, 234–235
ROLLUP, 236–237
GROUPING SETS, 235–236
GROUPING_ID, 235, 238–240
GZIP algorithm, 70

H

HAVING
CASE expressions, overview of, 52–54
GROUP BY and, 33
logical query processing order, 28
overview of, 35–36
HDInsight, 13
historical data. See temporal tables
HOLDLOCK, 332

I

IaaS (infrastructure as a service)
ABC flavors (Appliance, Box, Cloud), 13–14
IDENT_CURRENT, 255–257
identifier names
delimiting of, 30
quoted identifiers, 62–63
identity property, 254–257
IDENTITY_INSERT, 256–257
IF . . . ELSE flow element, 367–368
IIF, CASE expressions, 54
IMPLICIT_TRANSACTIONS, 319–320
IN
multivalued subquery examples, 136–139
precedence, 51–52

subqueries, NULL trouble and, 146–147
use of, 49
inconsistent analysis, 335, 344–345
INCREMENT BY, sequence objects and, 258–259
information schema views, metadata queries, 91
INFORMATION_SCHEMA, 91
infrastructure as a service (Iaas), 13–14
inline table-valued functions (TVFs)
exercises, 183–188
exercises, solutions for, 188–192
overview of, 171–172, 178–180
In-Memory OLTP, 18
inner joins, 107–109
exercises, 123–128
exercises, solutions for, 129–132
INSERT
DML triggers (data manipulation events), 387–388
exercises, 287–291
exercises, solutions for, 291–295
identity property and, 254–257
MERGE and, 274
OUTPUT and, 280–282
row-versioning isolation and, 339
table expressions, modifying data with, 276–278
views and, 172
INSERT EXEC, 251–252
INSERT SELECT, 251
INSERT VALUES, 249–251
INT, 51
integer sequences, cross joins and, 106–107
International Organization for Standardization (ISO), 2
INTERSECT
exercises, 204–208
exercises, solutions for, 208–211
precedence and, 201–202
use of, 196–199
INTERSECT (DISTINCT), 196–197
INTERSECT ALL, 196–199
INTO, SELECT INTO, 252–253
IS NOT NULL, 55–58
IS NULL
CASE expressions and, 54
overview of, 55–58
ISDATE, 88
ISO (International Organization for Standardization), 2
ISO/ANSI SQL-89 syntax
cross joins, 105
inner joins, 108–109
ISO/ANSI SQL-92 syntax
cross joins, 104
inner joins, 107–108
isolation, transactions, 320–322
deadlocks, 345–348
exercises, 348–359
overview of, 332–333
READ COMMITTED, 334–337
READ COMMITTED SNAPSHOT, 343–345
READ UNCOMMITTED, 333–334
REPEATABLE READ, 335–337
SERIALIZALBE, 337–338
SNAPSHOT, 339–343
SNAPSHOT, conflict detection, 341–343

J

JOIN
composite joins, 110–111
CROSS APPLY and, 180
cross-joins, 103–107
DELETE and, 265–266
derived tables, multiple references, 166
exercises, 123–128
exercises, solutions for, 129–132
inner joins, 107–109
multi-join queries, 113
natural joins, 111
non-equi joins, 111–112
OUTER APPLY and, 181
outer joins
COUNT aggregate, use of, 121–123
filtering attributes from nonpreserved side, 118–119
fundamentals of, 113–116
in multi-join queries, 119–121
missing values, inclusion of, 116–118
overview of, 103
self joins, 105–106
UPDATE and, 268–270
vs. subqueries, 137

K

keys, SQL queries and, 39–41

L

LAG, 378
offset window functions, 219–221
language independence, 2

language standards organizations, 2
LANGUAGE/DATEFORMAT, 75–76
LAST_VALUE, offset window functions, 219–221
LEAD, offset window functions, 219–221
LEFT, 65
LEN, 65–66
LIKE
  precedence, 51–52
  use of, 50, 71–73
linked history tables, 298–301
literals, date and time data, 74–78
local temporary tables, 374–375
LOCK_ESCALATION, 325
LOCK_TIMEOUT, 331–332
locks. See also isolation, transactions
  deadlocks, 337, 345–348
  exercises, 348–359
  lockable resource types, 324–325
  modes and compatibility, 322–324
  troubleshooting, 325–332
logical operators, 50
login, SQL Server databases, 17
loops, WHILE element and, 368–370
LOWER, 68–69
LTRIM, 69

M

massively parallel processing (MPP), 13
master database, SQL Server architecture, 16–18
MAX, 33–35, 61
MAXVALUE, sequence objects and, 258–259
MERGE
  DML triggers (data manipulation events), 387–388
  exercises, 287–291
  exercises, solutions for, 291–295
  merging data overview, 271–275
  OUTPUT and, 284–285
  table expressions, modifying data with, 276–278
metadata, querying of, 89–92
Microsoft Analytics Platform System (APS), 13–14
Microsoft Azure SQL Data Warehouse, 14
Microsoft Azure SQL Database. See Azure SQL Database
Microsoft SQL Server. See SQL Server
MIN, 33–35
MINVALUE, sequence objects and, 258–259
missing values, 6–7
  in outer joins, 116–118
model database, SQL Server architecture, 16–18
MONTH, 79–80, 87–88
msdb database, SQL Server architecture, 16–18
MTD, aggregate window functions, 223–224
multi-join queries
  exercises, 123–128
  exercises, solutions for, 129–132
  outer joins in, 119–121
  use of, 113
multiplication (*), 51–52
multiset, 3
multi-statement table-valued function (TVF), 178
multivalued subqueries
  defined, 133
  self-contained subquery examples, 136–139

N

N (National), data type prefix, 50, 60–61
names, schemas and objects, 18–19. See also aliases
natural joins, 111
NCHAR, 50, 60–61
negative (-), 51–52
nested DML, 285–287
nested queries, 133
  derived tables and, 165–166
NEWID, 251, 383
NEXT VALUE FOR, 259–262
next values, subquery returns, 143–144
no action, foreign-key constraints, 24
NOLOCK, 332
non-equi joins. See equi and non-equi joins
nonkey attributes, normalization forms, 8–10
nonrepeatable reads, 335, 344–345
normal forms. See normalization
normalization, overview, 7–10
NOT
  combining logical expressions, 50
  multivalued subqueries and, 137
  precedence, 51–52
  subqueries, NULL trouble and, 146–147
NOT EXISTS, NULL trouble and, 147
NOT IN, NULL trouble and, 146–147
NOT NULL, grouping and, 238–240
NTILE, window function ranking, 216–219
NULL
  aggregate functions and, 34–35
  CASE expressions and, 226
  data integrity and, 22–25

DISTINCT and duplicates, 194
GROUPING and GROUPING_ID functions and, 238–240
IF . . . ELSE flow element, 367–368
INSERT VALUES and, 250
INTERSECT and, 197
missing values, overview, 6–7
multivalued subqueries and, 137
outer joins, filtering attributes, 118–119
outer joins, fundamentals of, 115–116
outer joins, in multi-join queries, 119–121
overview of, 54–58
scalar subqueries and, 135
SET NULL and, 24
string concatenation and, 63–65
subqueries, troubles with, 145–147
tables, creating, 21
variable values and, 363
NUMERIC, 51
NVARCHAR, 50, 60–61

O

object names, schema qualifying of, 29
object, defined, 3
OBJECT_ID, table creation, 20
OBJECTPROPERTY, 92
objects, SQL Server architecture, 18–19. See also programmable objects
offset window functions, 219–221
OFFSET-FETCH
data modification with, 278–280
derived tables and, 162
exercises, 287–291
exercises, solutions for, 291–295
overview of, 46–47
views and ORDER BY clause, 173–174
OLTP (online transactional processing) databases, 10–12
ON
outer joins, fundamentals of, 115–116
outer joins, in multi-join queries, 119–121
ON DELETE CASCADE, 24
online transactional processing (OLTP) databases, overview, 10–12
operators. See also APPLY; also JOIN; also PIVOT; also UNPIVOT
compound operators, UPDATE and, 267–268
overview of, 49–52
precedence rules, 51–52
string concatenation, 63–65
OR
combining logical expressions, 50
precedence, 51–52
predicates and, 49
ORDER BY
CASE expressions, overview of, 52–54
circumventing unsupported logical phrases, 203–204
derived tables and, 162
GROUP BY and, 33
INTERSECT ALL and, 197–199
logical query processing order, 28
offset window functions and, 220–221
OFFSET-FETCH, overview of, 46–47
overview of, 41–43
set operators and, 193
TOP filters and, 45–46
views, table expressions, 172–174
window functions, overview of, 47–49, 214–215
ordinal positions, ORDER BY clause and, 43
OUTER APPLY
JOIN and, 181
overview of, 180–183
outer joins
COUNT aggregate, use of, 121–123
exercises, 123–128
exercises, solutions for, 129–132
filtering attributes from nonpreserved side, 118–119
fundamentals of, 113–116
in multi-join queries, 119–121
missing values, inclusion of, 116–118
outer queries, 133. See also subqueries; also table expressions
OUTPUT
DELETE and, 282–283
exercises, 287–291
exercises, solutions for, 291–295
INSERT and, 280–282
MERGE and, 284–285
nested DML, 285–287
overview of, 280
UPDATE and, 283–284
OVER
empty parens () and, 221–222
sequence objects and, 260–262
window functions, overview of, 47–49, 213–216

P

PaaS (platform as a service)
ABC flavors (Appliance, Box, Cloud), 13–14
Parallel Data Warehouse (PDW), 13
parentheses, (), 51–52
PARSE
date and time functions, 81–83
date and time literals, 77–78
parsing, batches and, 364
PARTITION BY
INTERSECT ALL and, 197–199
window functions, overview of, 47–49, 215
partitions, offset window functions, 219–221
PATINDEX, 66
percent (%) wildcard, 71–72
PERCENT, TOP filters and, 45
performance
dynamic SQL and, 379–383
object names, schema qualifying, 29
query filters and, 32
row-versioning isolation and, 339
sp_executesql stored procedure, 380–381
SQL Server architecture, physical layout, 17–18
stored procedures and, 385
PERIOD FOR SYSTEM_TIME, temporal table creation, 298–301
permissions, database schemas and objects, 18–19
phantom reads, 337–338
SNAPSHOT and, 339
PIVOT
dynamic SQL and, 379, 381–383
pivoting data, overview, 227–229
pivoting data
exercises, 241–244
exercises, solutions for, 245–248
overview of, 224–226
unpivoting data, 230–234
with PIVOT operator, 227–229
within grouped queries, 226–227
platform as a service (PaaS), 13–14
point-in-time analysis. See temporal tables
policy enforcement, DDL triggers (data definition events), 388–390
PolyBase, 13
positive (+), 51–52
predicates. See also specific predicate names
LIKE predicate, 71–73
NULLs, overview of, 54–58
overview of, 49–52
predicate logic, 4
relational model, overview, 5–6
previous table versions. See temporal tables
previous values, subquery returns, 143–144
primary-key constraints, data integrity, 22
private cloud, ABC flavors (Appliance, Box, Cloud), 13–14
procedural data integrity, 22–25
programmable objects
batches and, 363–367
cursors, 370–373
dynamic SQL, 379–383
error handling, 390–394
flow element, 367–370
routines, 383
stored procedures, 385–386
triggers, 387–390
user-defined functions, 383–384
temporary tables, 374–379
variables, 361–363
propositions, relational model overview, 5–6
public cloud, ABC flavors (Appliance, Box, Cloud), 13–14

Q

queries. See also SELECT statement; also subqueries; also table expressions
aliases, use of, 36–41
all-at-once operations, 58–59
CASE expressions, overview of, 52–54
character data, working with, 60–63
date and time data, working with, 73–89
exercises, 93–97
exercises, solutions for, 97–102
FROM clauses, overview of, 29–30
GROUP BY clause, overview of, 32–35
HAVING clause, overview of, 35–36
identifier names, delimiting of, 30
in temporal tables, 304–310
LIKE predicate, 71–73
logical query processing, overview, 27–29
nested and outer queries, defined, 133
NULLs, overview of, 54–58
of metadata, 89–92
OFFSET-FETCH filter, overview of, 46–47
operators and functions, overview of, 63–73
ORDER BY clause, overview of, 41–43
predicates and operators, overview of, 49–52
query clause, defined, 29

query phrase, defined, 29
scalar subqueries, 133
SELECT clause, overview of, 36–41
TOP filters, overview of, 44–46
WHERE clause, overview of, 31–32
window functions, overview of, 47–49
quoted identifiers, 62–63

R

RAND, 383
RANK, window function ranking, 216–219
RDBMSs (relational database management systems)
ABC flavors (Appliance, Box, Cloud), 12–14
defined, 1
language independence, 2
relational model, overview, 4–10
READ COMMITTED
conflict detection, SNAPSHOT, 341–343
default isolation levels, 323
isolation, overview of, 332–337, 345
SNAPSHOT and, 340, 343
READ COMMITTED LOCKS, 343
READ COMMITTED SNAPSHOT
conflict detection, SNAPSHOT, 341–343
default isolation levels, 323, 332–333
isolation, overview, 339, 343–345
READ UNCOMMITTED, isolation overview, 332–334, 345
referenced table, foreign-key constraints, 23–24
referencing table, foreign-key constraints, 23–24
relation variable, use of term, 5
relational database management systems (RDBMSs)
ABC flavors (Appliance, Box, Cloud), 12–14
defined, 1
language independence, 2
relational model, overview, 4–10
relational model
constraints, 7
missing values, 6–7
normalization, 7–10
overview of, 4–10
propositions, predicates, and relations, 5–6
REPEATABLE READ, 332–333, 335–337, 345
REPLACE, 66–67
REPLICATE, 67–68
resolution, batches and, 366
Resource database, SQL Server architecture, 16
resource types, locking of, 324–325
RESTART WITH, sequence objects and, 259
restoring tables. See temporal tables
RIGHT, 65
ROLLBACK TRAN (TRANSACTION)
lock modes and compatibility, 322–324
overview of, 319–322
triggers and, 387
ROLLUP, 235–237
routines
overview of, 383
stored procedures, 385–386
triggers, 387–390
user-defined functions, 383–384
row constructors, 270
row versioning
isolation and, 320–321
READ COMMITTED SNAPSHOT, 343–345
SNAPSHOT, 339–343
ROW_NUMBER
EXCEPT ALL and, 200–201
INTERSECT ALL and, 197–199
window functions, overview of, 47–49
window functions, ranking of, 216–219
rows, 4
ROWS BETWEEN
offset window functions, 220–221
window functions, overview of, 215
RTRIM, 69
running aggregates, subqueries and, 144–145

S

scalar subqueries
defined, 133
scalar variables and, 361–363
self-contained subquery examples, 134–135
scalar UDFs (user-defined functions), 383–384
scalar variables, 361–363
SCHEMABINDING, views and, 176–177
schemas
dbo-schema, 20
identifier names, 30
information schema views, metadata queries, 91
object names, qualifying of, 29
SQL Server architecture, 18–19
SCOPE_IDENTITY, 255–257
searched CASE expressions, 52–54
security
database schemas and objects, SQL Server, 18–19
stored procedures, 385

SELECT
INSERT SELECT, 251
pivoting with grouped queries, 226–227
variables and, 362–363
window functions, overview of, 215–216
window functions, ranking of, 218–219
SELECT * FROM, views and, 172
SELECT *, views and, 172
SELECT INTO, 252–253
SELECT statement. See also table expressions
all-at-once operations, 58–59
CASE expressions, overview of, 52–54
character data, working with, 60–63
date and time data, working with, 73–89
derived tables, assigning aliases, 163–164
elements of, 27–29
exercises, 93–97
exercises, solutions for, 97–102
FROM, overview, 29–30
GROUP BY, overview of, 32–35
HAVING, overview of, 35–36
LIKE predicate, 71–73
logical query processing order, 28
metadata, querying of, 89–92
NULLs, overview of, 54–58
OFFSET-FETCH filter, overview of, 46–47
operators and functions, overview of, 63–71
ORDER BY, overview of, 41–43
predicates and operators, overview of, 49–52
SELECT, overview of, 36–41
TOP filters, overview of, 44–46
views and, 172
WHERE, overview of, 31–32
window functions, overview of, 47–49
self joins, 105–106
self-contained subqueries
defined, 133
multivalued subquery examples, 136–139
scalar subquery examples, 134–135
semicolon, use in statements, 29
SEQUEL (Structured English QUEry Language), 2
sequence object, 258–262
SERIALIZALBE
isolation, overview, 332–333, 337–338, 345
SNAPSHOT and, 340
SERVERPROPERTY, 92
session ID, troubleshooting blocks, 326–332
SET
quoted identifiers, 62–63
UPDATE and, 267–268
variables and, 361–363
set attributes
normalization forms, 7–10
propositions, predicates, and relations, 5–6
use of term, 5
SET DEFAULT, 24
SET NOCOUNT ON, 386
SET NULL, 24
set operators
EXCEPT, 199–201
exercises, 204–208
exercises, solutions for, 208–211
INTERSECT, 196–199
overview of, 193–194
precedence and, 201–202
UNION, 194–196
unsupported logical phrases, circumventing, 203–204
set theory
overview of, 3
predicate logic and, 4
simple CASE expressions, 52–54
skipping option, OFFSET-FETCH, 46–47
SMALLDATETIME
data types, 73–74
literals, 74–78
working with date and time separately, 78–79
SNAPSHOT
conflict detection, 341–343
isolation, overview of, 339–343, 345
READ COMMITTED SNAPSHOT, 343–345
snowflake dimension, 11
solutions. See exercises
sorting, ORDER BY clause and, 41–43
sp_columns, 92
sp_executesql stored procedure, 380–381
sp_help, 91–92
sp_helpconstraint, 92
sp_tables, 91
SQL (Structured Query Language), overview of, 1–3
database system types, 10–12
history and use of, 2–3
predicate logic, 4
relational model, overview of, 4–10
constraints, 7
missing values, 6–7
normalization, 7–10
propositions, predicates, and relations, 5–6
relational models, language independence, 2
set theory, overview of, 3

standards for use, 2
statement termination, semicolons, 29
SQL Data Warehouse, 14
SQL Server
architecture overview, 12–19
ABC flavors (Appliance, Box, Cloud), 12–14
databases, 15–18
instances, 14–15
physical layout, 17–18
schemas and objects, 18–19
box product installation, 396–401
database engine installation, 396–401
file extensions, 18
operator precedence rules, 51–52
source code, downloading of, 402–404
SQL Server Books Online, 410–413
unique index, 22–23
SQL Server authenticated login, 17
SQL Server Management Studio (SSMS)
download and installation of, 402
getting started with, 404–410
session ID, troubleshooting blocks, 326–332
temporal tables, creating, 299
SQL_VARIANT, 262
star schema, overview of, 11–12
START WITH, sequence objects and, 258–259
storage, character data and, 60–61
stored procedures, 385–386
triggers, 387–390
string concatenation, 63–65
STRING_SPLIT, 71
Structured Query Language (SQL), overview of, 1–3
STUFF, 68
subqueries. See also table expressions
column names, substitution errors, 147–149
correlated subqueries
EXISTS and, 141–143
overview of, 139–141
exercises, 150–154
exercises, solutions for, 154–159
multivalued subqueries, defined, 133
NULL trouble, 145–147
overview of, 133
previous or next values, returning, 143–144
running aggregates and, 144–145
scalar subqueries, defined, 133
self-contained
defined, 133
multivalued subquery examples, 136–139
scalar subquery examples, 134–135
table subqueries, defined, 133
vs. joins, 137
window functions, overview of, 214
SUBSTRING, 65
subtraction (-), 51–52
SUM, 33–35. See also window functions
surrogate keys, identity property and, 254–257
SWITCHOFFSET, 83
syntax
batch as unit of parsing, 364
cross joins, 104–105
inner joins, 107–109
sys.dm_as_waiting_tasks, 330–331
sys.dm_exec_connections, 328
sys.dm_exec_input_buffer, 329
sys.dm_exec_requests, 330–331
sys.dm_exec_sessions, 329
sys.dm_exec_sql_text, 328
sys.dm_tran_locks, 328
SYSDATETIME
current date and time functions, 80–81
default constraints, defining, 25
INSERT SELECT and, 251
SYSDATETIMEOFFSET, 80–81
system versioned temporal tables. See temporal tables
SYSTEM_TIME, temporal table queries, 304–310
SYSTEM_VERSIONING, temporal table creation, 298–301
SYSUTCDATETIME, 80–81

T

table expressions
APPLY operator, 180–183
common table expressions (CTEs)
arguments, using, 167–168
column aliases, assigning, 167–168
defining multiple CTEs, 168–169
multiple references in, 169
overview, 167
recursive CTEs, 169–171
derived tables
arguments, use of, 164–165
column aliases, assigning, 163–164
multiple references, 166
nesting, 165–166
overview of, 161–162
exercises, 183–188, 287–291
exercises, solutions for, 188–192, 291–295

table expressions *continued*
inline table-valued functions (TVFs), overview of, 171–172, 178–180
modifying data with, 276–278
overview of, 161
views
CHECK OPTION, 177–178
ENCRYPTION, 175–176
ORDER BY, 172–174
overview of, 171–172
SCHEMABINDING, 176–177
table subqueries, defined, 133. See also table expressions
tables
check constraints, defining, 24–25
creating tables, 19–22
defining data integrity, 22–25
foreign-key constraints, defining, 23–24
identifier names, 30
primary-key constraints, defining, 23
unique key constraints, defining, 23
table-valued functions (TVFs), inline, 171–172
table-valued parameters (TVPs), 378–379
table-valued UDFs (user-defined functions), 383–384
tempdb, SQL Server
architecture overview, 16–18
local temporary tables, 374–375
SNAPSHOT, 339–343
temporary table variables, 377–378
temporal tables
creating, 298–301
exercises, 311–313
exercises, solutions for, 313–317
modifying data, 301–304
overview of, 297
querying data, 304–310
temporary tables
global temporary tables, 376–377
local temporary tables, 374–375
overview of, 374
table types, 378–379
table variables, 377–378
THEN, CASE expressions, 53–54
tiebreakers, ORDER BY clause and, 45–46
TIME
data types, 73–74
literals, 74–78
working with date and time separately, 78–79
time data
aggregate window functions, 223–224
AT TIME ZONE, 84–85
CAST, CONVERT, and PARSE, 81–83
current date and time functions, 80–81
DATEADD, 85
DATEDIFF and DATEDIFF_BIG, 86–87
DATENAME, 88
DATEPART, 87
EOMONTH, 89
FROMPARTS, 88–89
ISDATE, 88
literals, 74–78
missing values, in outer joins, 116–118
SWITCHOFFSET, 83
types of, 73–74
working with date and time separately, 78–79
YEAR, MONTH, and DAY, 87–88
TODATETIMEOFFSET, 83–84
TOP (100) PERCENT, views and, 173–174
TOP filters
data modification with, 278–280
derived tables and, 162
exercises, 287–291
exercises, solutions for, 291–295
overview of, 44–46
views and ORDER BY clause, 173–174
transaction log, durability and, 321
transactions
deadlocks, 345–348
exercises, 348–359
isolation
overview of, 332–333
READ COMMITTED, 334–337
READ COMMITTED SNAPSHOT, 343–345
READ UNCOMMITTED, 333–334
REPEATABLE READ, 335–337
SERIALIZALBE, 337–338
SNAPSHOT, 339–343
SNAPSHOT, conflict detection, 341–343
summary of, 345
locks and blocking, 322–325
locks and blocking, troubleshooting, 325–332
overview of, 319–322
vs. batches, 363
triggers, 387–390
troubleshooting
locks and blocking, 325–332
OUTPUT, use of, 280
table expressions, modifying data with, 276–278
TRUE
check constraints and, 24–25
IF . . . ELSE flow element, 367–368

meaning of, 32
NULLs, overview of, 54–58
WHILE flow element, 368–370
TRUNCATE, 262–266
TRY . . . CATCH, error handling, 390–394
TRY_, date and time functions, 81–83
T-SQL
coding style, 21–22
database system types, 10–12
language independence, 2
predicate logic, 4
relational model, overview of, 4–10
constraints, 7
missing values, 6–7
normalization, 7–10
propositions, predicates, and relations, 5–6
set theory, 3–4
SQL standards and, 3
SQL, background information, 1–3
whole, use of, 3
tuples, 4–5
type; propositions, predicates, and relations, 6

U

UDF (user-defined functions), routines, 383–384
uncommitted reads, 333–334
underscore (_) wildcard, 72
Unicode data types, 50
character data, 60–61
UNION
exercises, 204–208
exercises, solutions for, 208–211
precedence and, 201–202
use of, 194–196
UNION (DISTINCT), 194–196
UNION ALL, 194–195
unique constraint, data integrity, 23
unique index, SQL Server
primary-key constraints and, 22
unique constraints, defining, 23
UNIQUE, NULLs and, 58
UNKNOWN
check constraints and, 24–25
IF . . . ELSE flow element, 367–368
meaning of, 32
NULLs, overview of, 54–58
outer joins, filtering attributes, 118–119
outer joins, in multi-join queries, 119–121
scalar subqueries and, 135
subqueries, NULL trouble and, 146–147
WHILE flow element, 368–370
unknown, missing values, 6–7
UNPIVOT, 233–234
unpivoting data, 230–234
exercises, 241–244
with APPLY, 231–233
with UNPIVOT, 233–234
UPDATE
assignment UPDATE, 270–271
based on a join, 268–270
DML triggers (data manipulation events), 387–388
exercises, 287–291
exercises, solutions for, 291–295
MERGE and, 274
OUTPUT and, 283–284
overview of, 266–268
row-versioning isolation and, 339
table expressions, modifying data with, 276–278
views and, 172
UPPER, 68–69
USE, table creation, 20
user-defined functions (UDF), routines, 383–384

V

VALUES, 250–251
VAR element, 60–61
VARCHAR, 21, 50, 60–61
variables, as programmable objects
batches and, 365
overview of, 361–363
temporary table variables, 377–378
vector expressions, 270
views, table expressions
CHECK OPTION, 177–178
ENCRYPTION, 175–176
exercises, 183–188
exercises, solutions for, 188–192
ORDER BY and, 172–174
overview of, 171–172
SCHEMABINDING, 176–177
virtual machine (VM), ABC flavors (Appliance, Box, Cloud), 13–14

W

WHEN MATCHED THEN, 273–275
WHEN NOT MATCHED THEN, 273–275

WHEN, CASE expressions, 53–54
WHERE
  CASE expressions, overview of, 52–54
  derived tables, arguments, 164–165
  derived tables, assigning aliases, 163
  logical query processing order, 28
  outer joins, filtering attributes, 118–119
  outer joins, fundamentals of, 115–116
  overview of, 31–32
  UPDATE and, 267–268
  UPDATE based on a join, 268–270
WHILE flow element, 368–370
whole, defined, 3
wildcards
  % (percent) wildcard, 71–72
  _ (underscore) wildcard, 72
  [^character list or range] wildcard, 73
  [character-character] wildcard, 72–73
  [list of characters] wildcard, 72
  LIKE predicate and, 71–73
window functions
  aggregate window functions, 221–224
  exercises, 241–244
  exercises, solutions for, 245–248
  offset functions, 219–221
  overview of, 47–49, 213–216
  ranking of, 216–219
Windows authenticated login, SQL Server databases, 17
WITH NOCHECK, 25
WITH TIES, ORDER BY clause and, 45–46
WITH, common table expressions (CTEs), 167

X

XACT_ABORT, 320

Y

YEAR, 87–88
  filtering date ranges, 79–80
YTD, aggregate window functions, 223–224